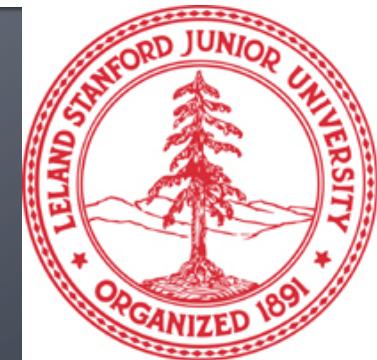


Community Detection in Graphs: Finding overlaps

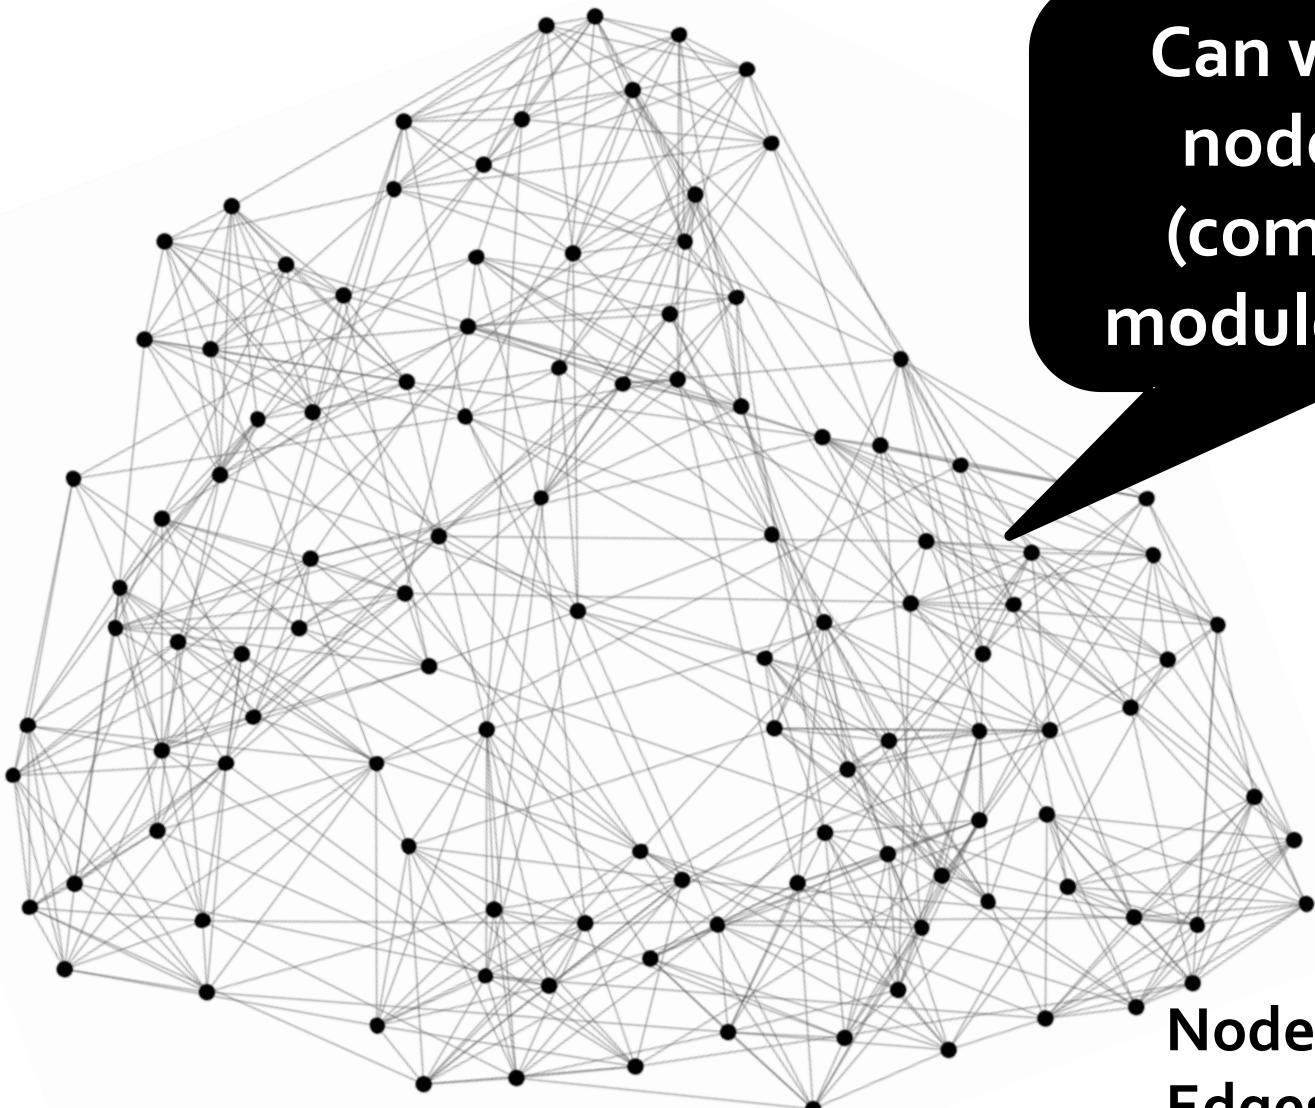
CS246: Mining Massive Datasets

Jure Leskovec, Stanford University

<http://cs246.stanford.edu>



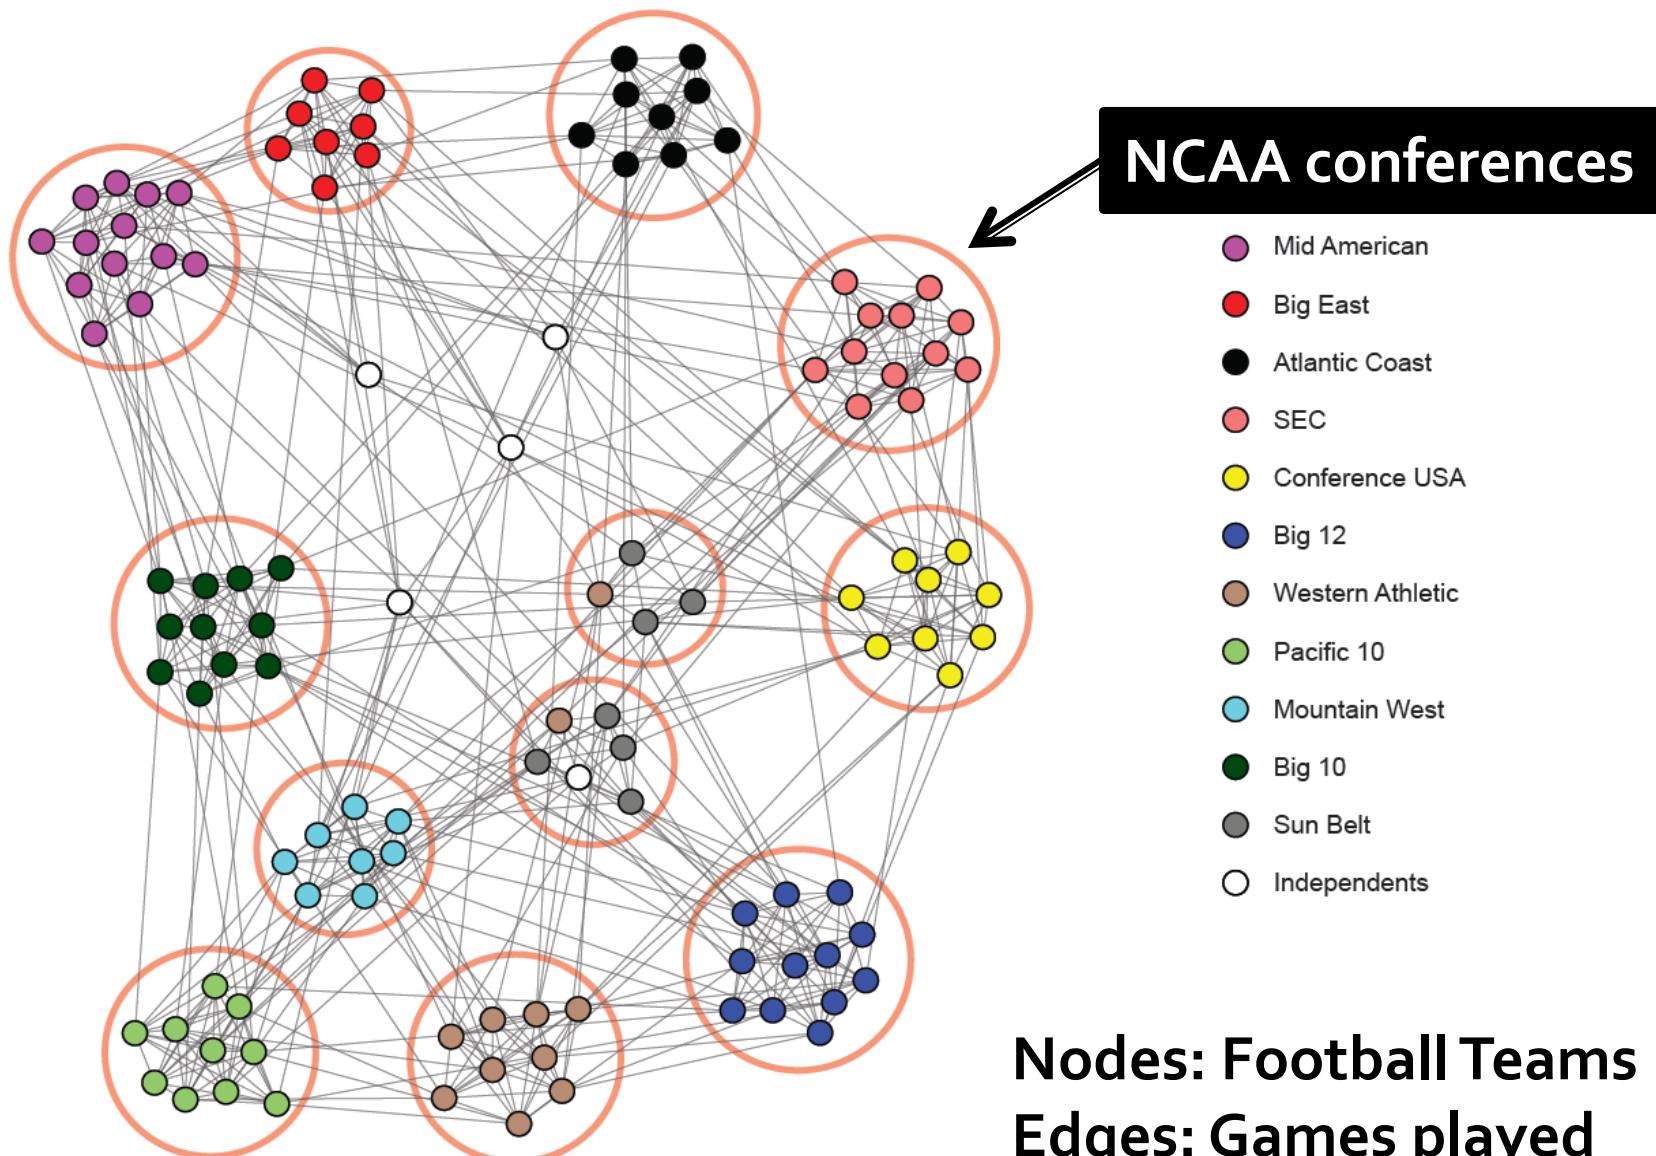
Identifying Communities



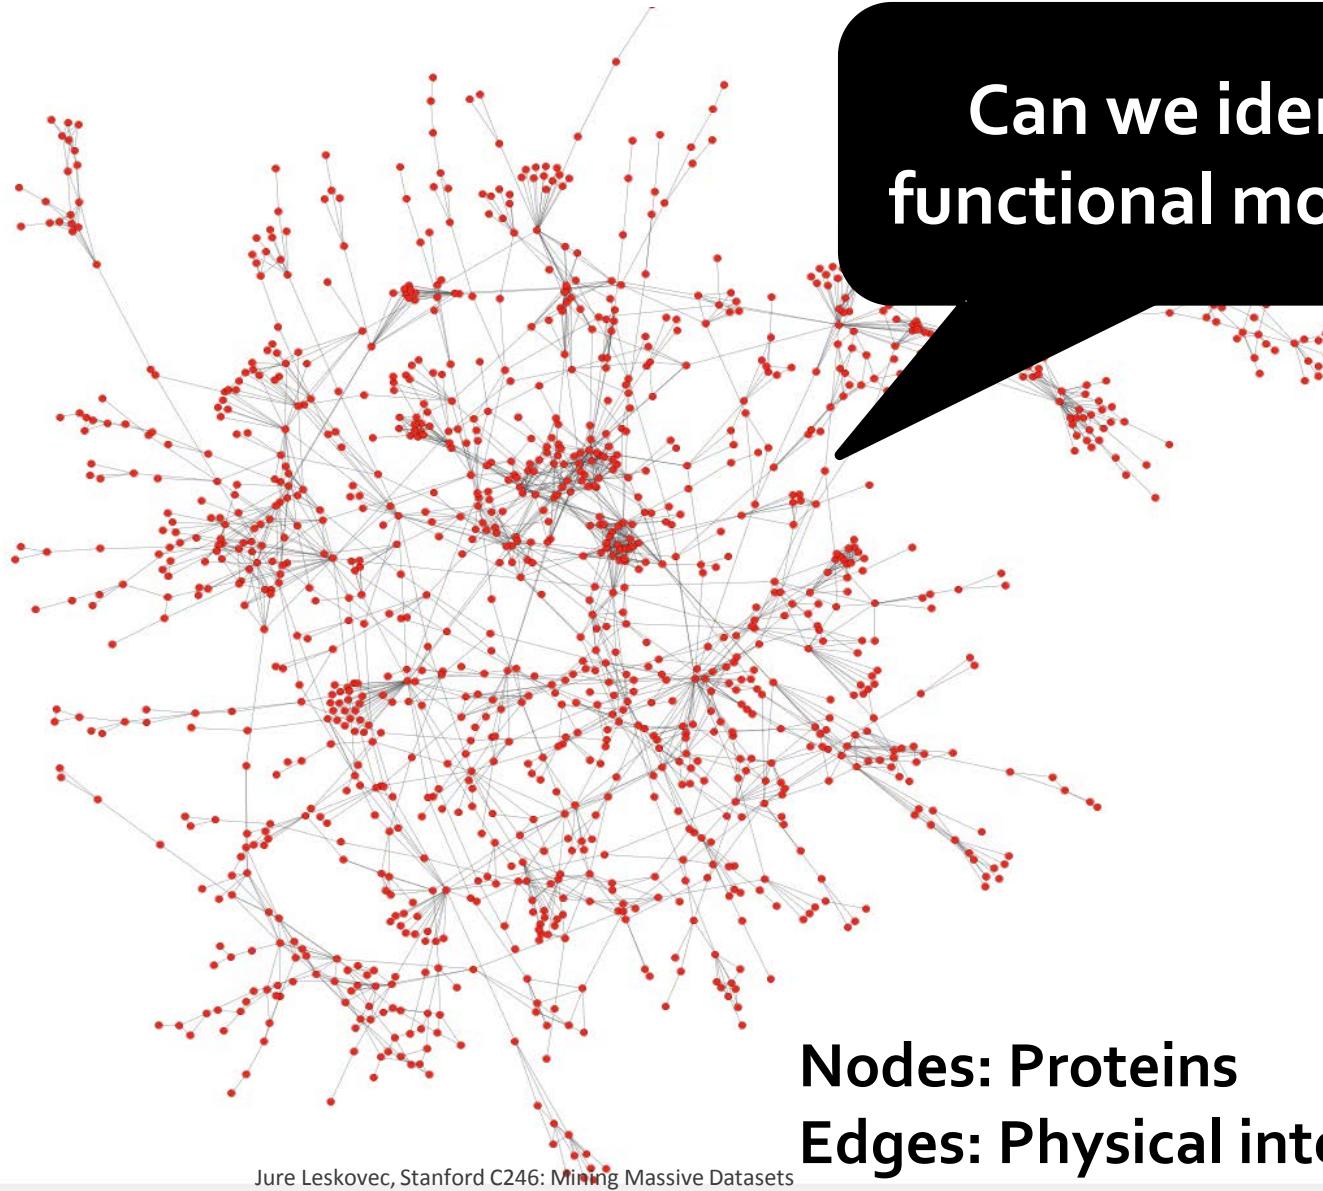
Can we identify
node groups?
(communities,
modules, clusters)

Nodes: Football Teams
Edges: Games played

NCAA Football Network

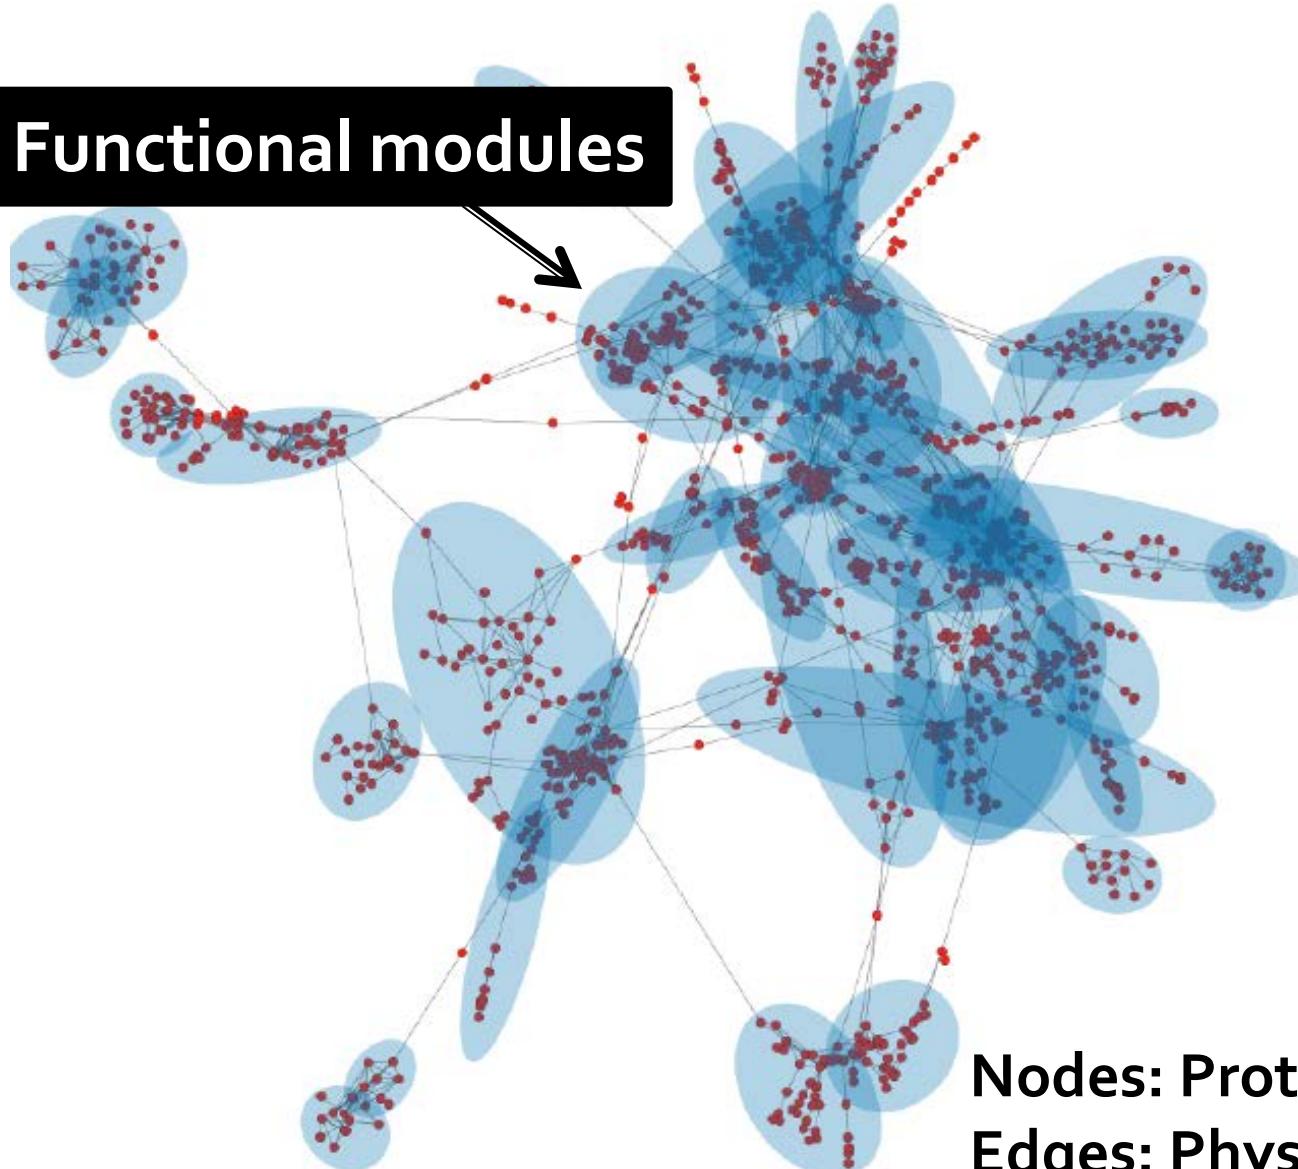


Protein-Protein Interactions

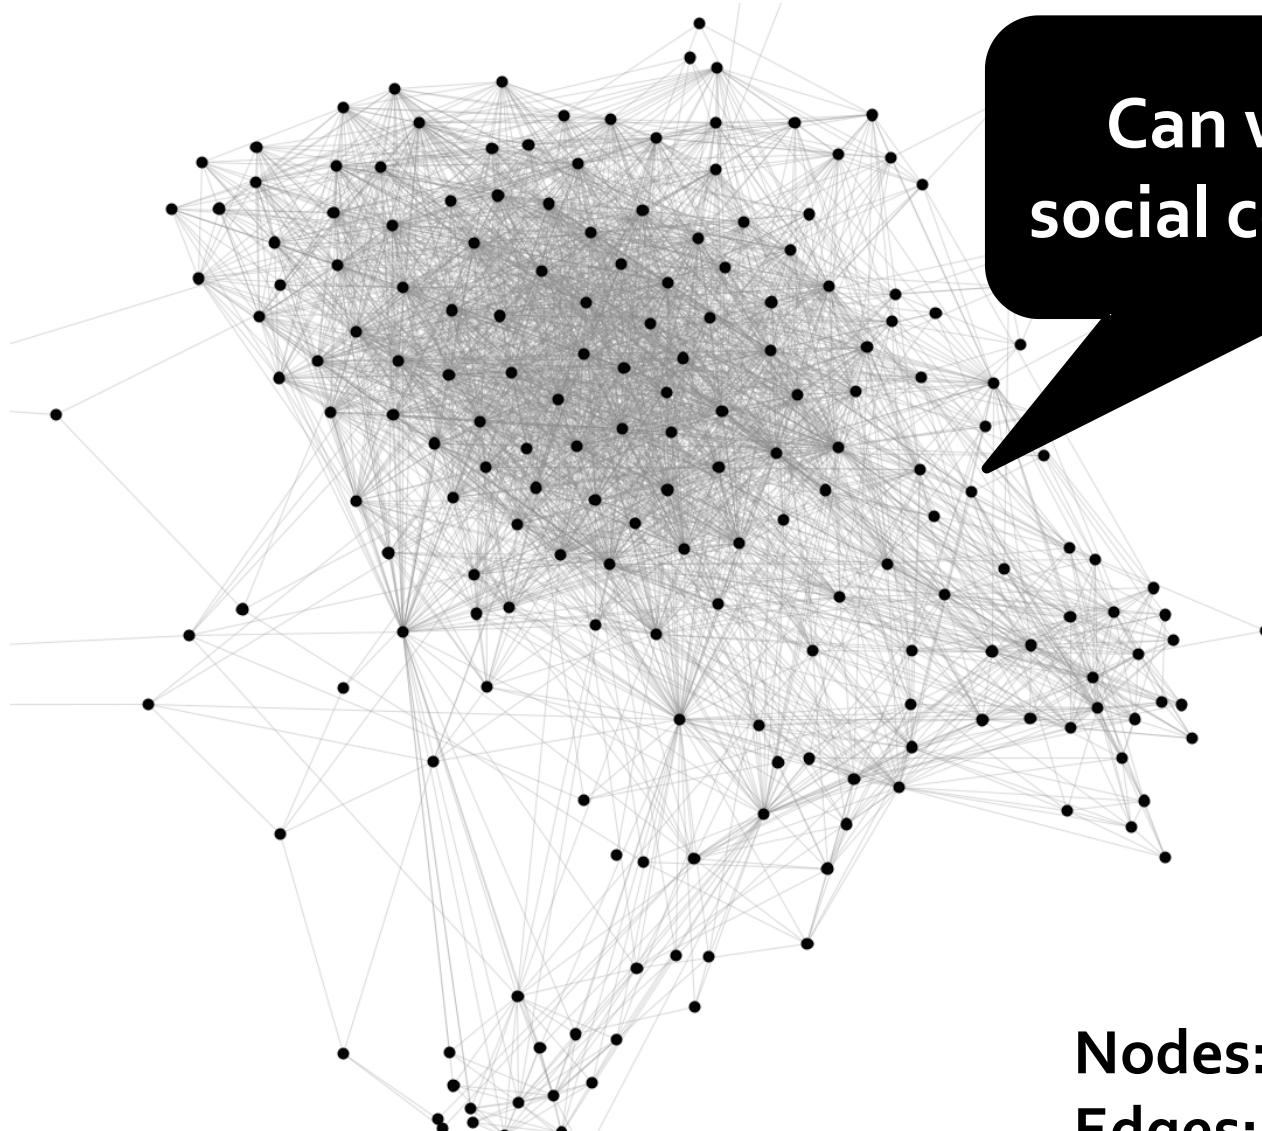


Protein-Protein Interactions

Functional modules



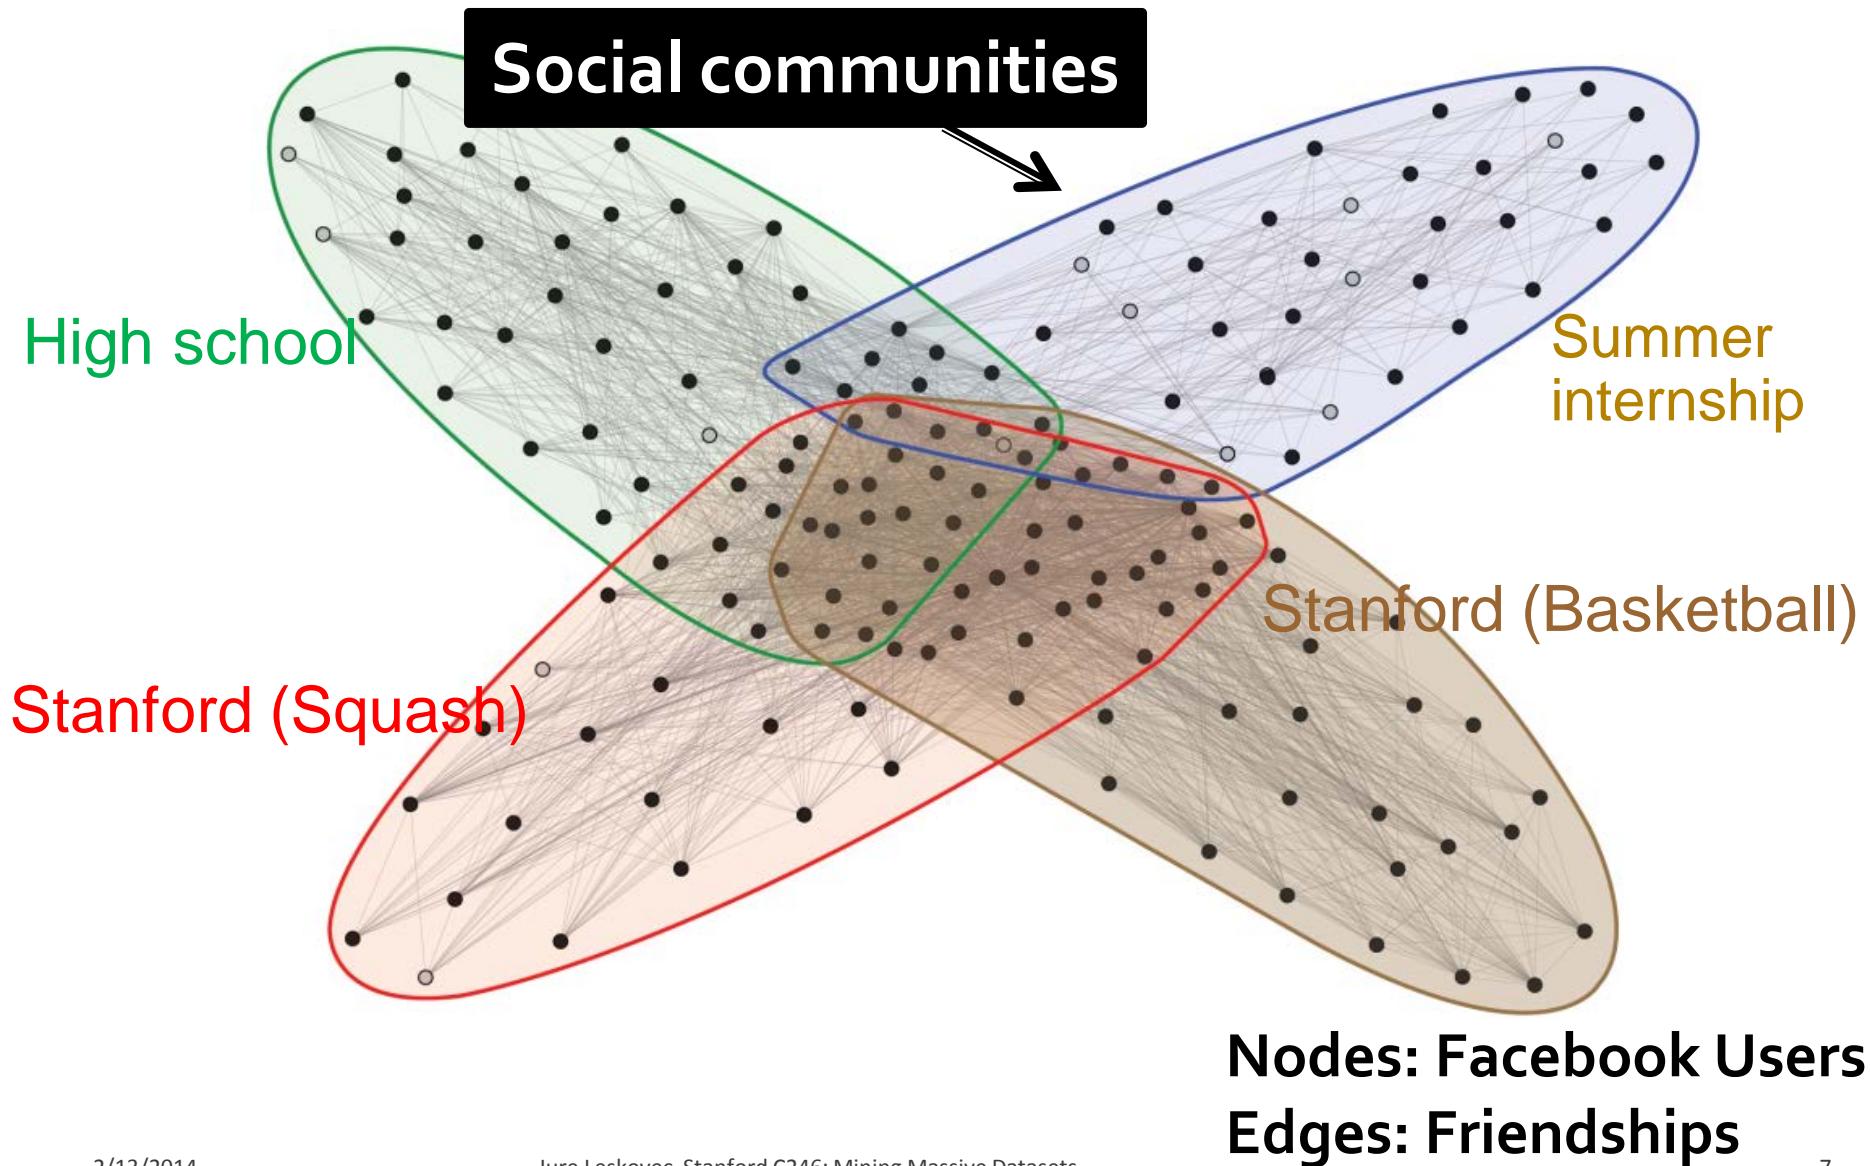
Facebook Network



Can we identify
social communities?

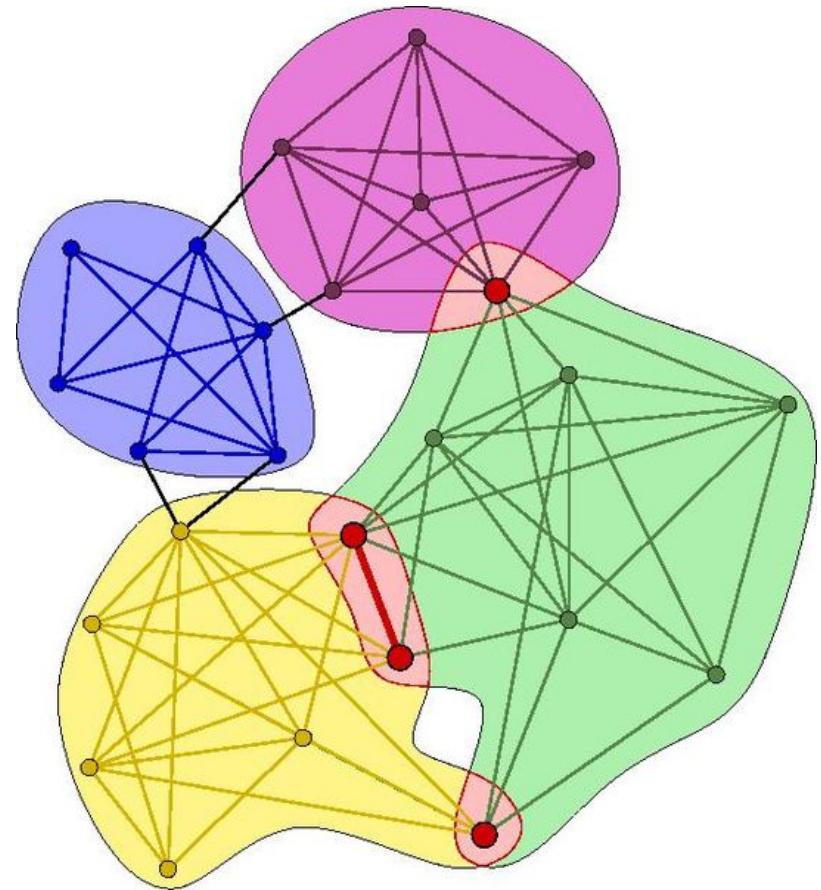
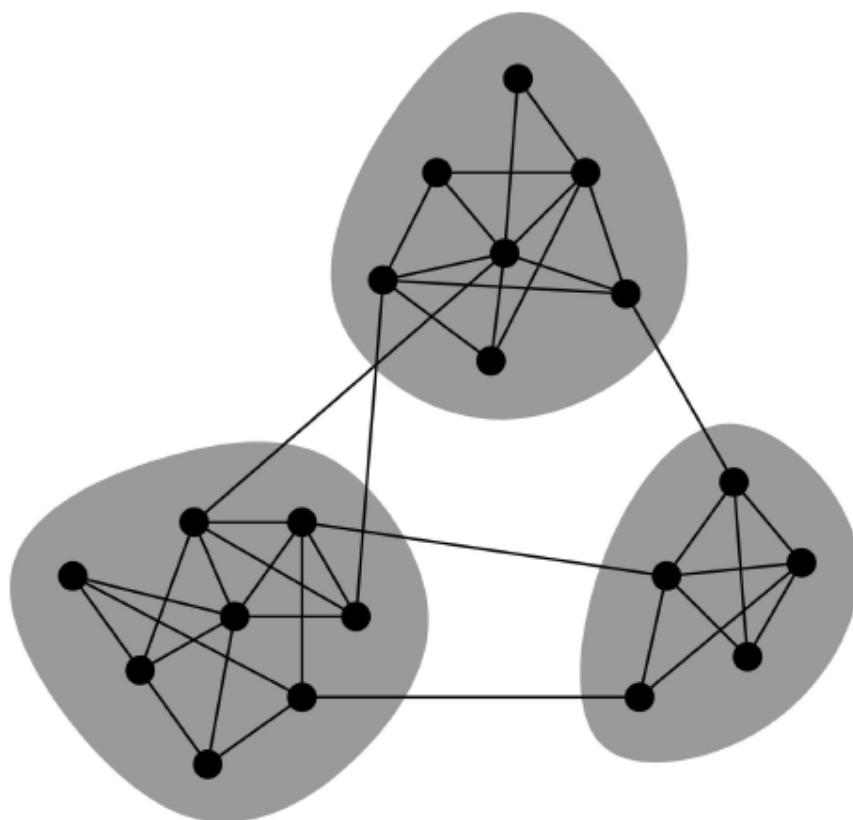
Nodes: Facebook Users
Edges: Friendships

Facebook Network

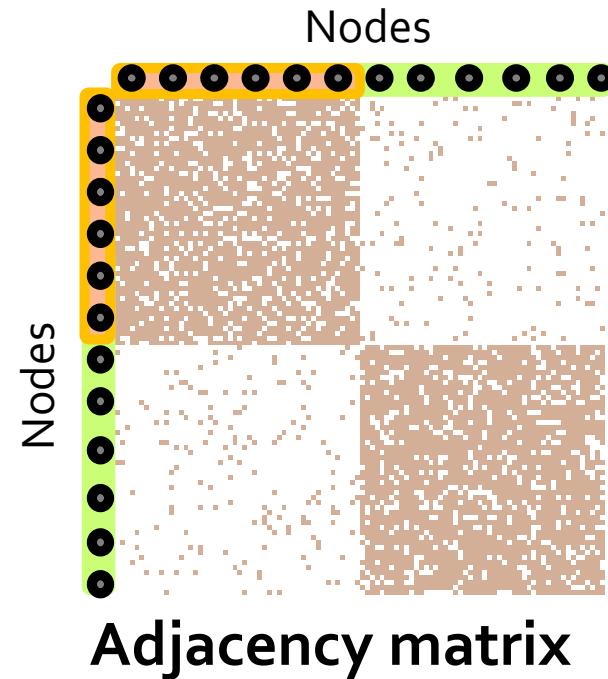
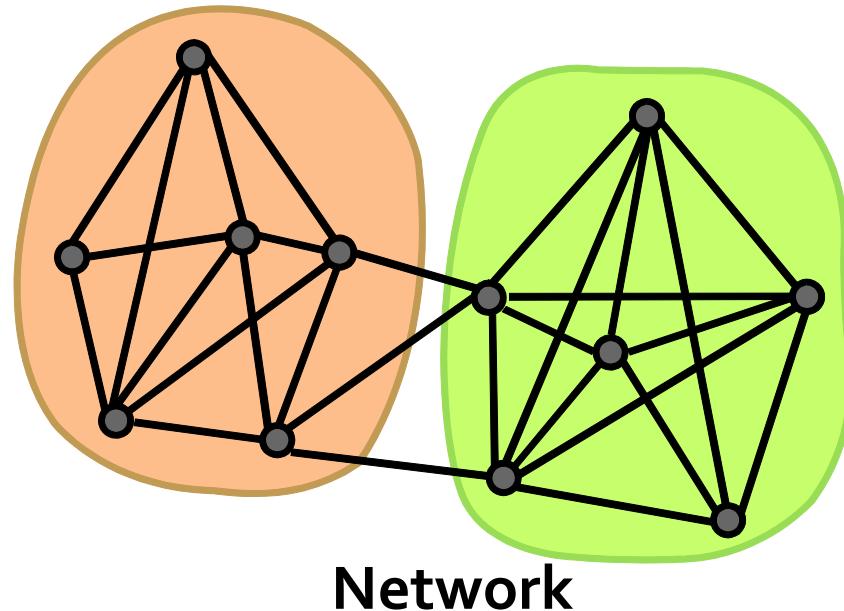


Overlapping Communities

- Non-overlapping vs. overlapping communities

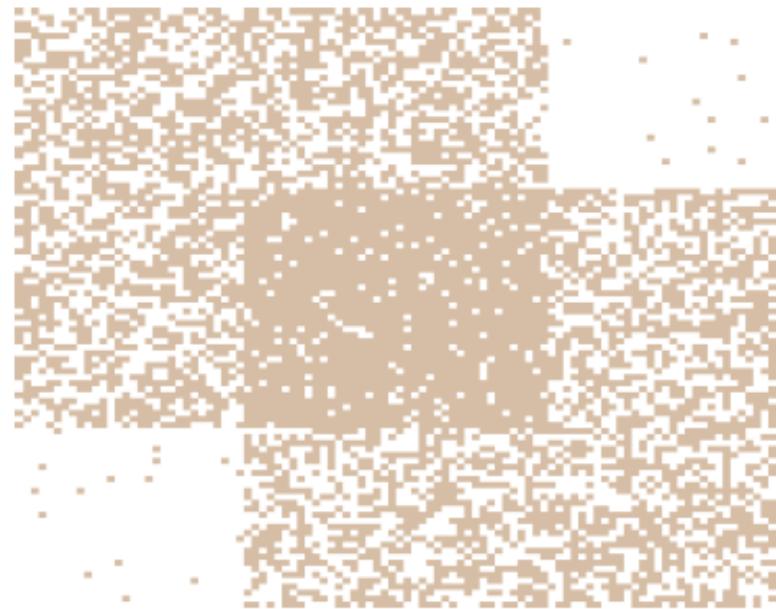
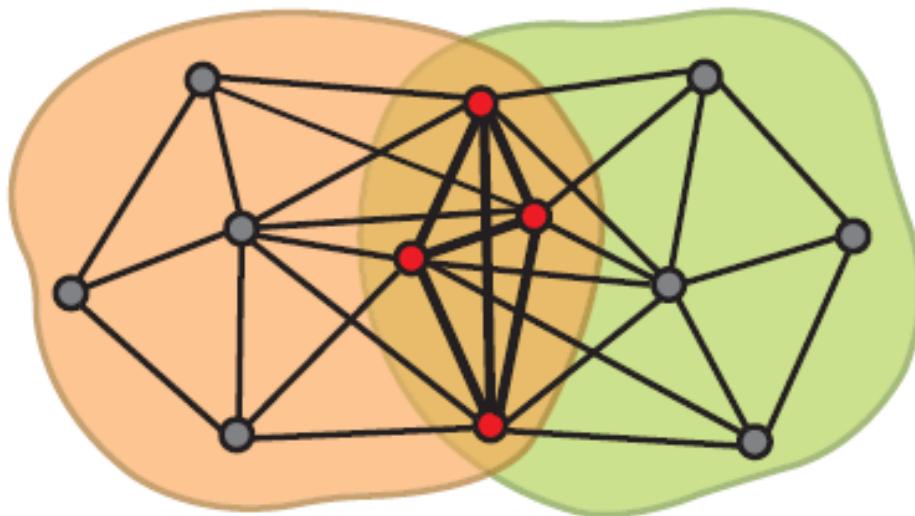


Non-overlapping Communities



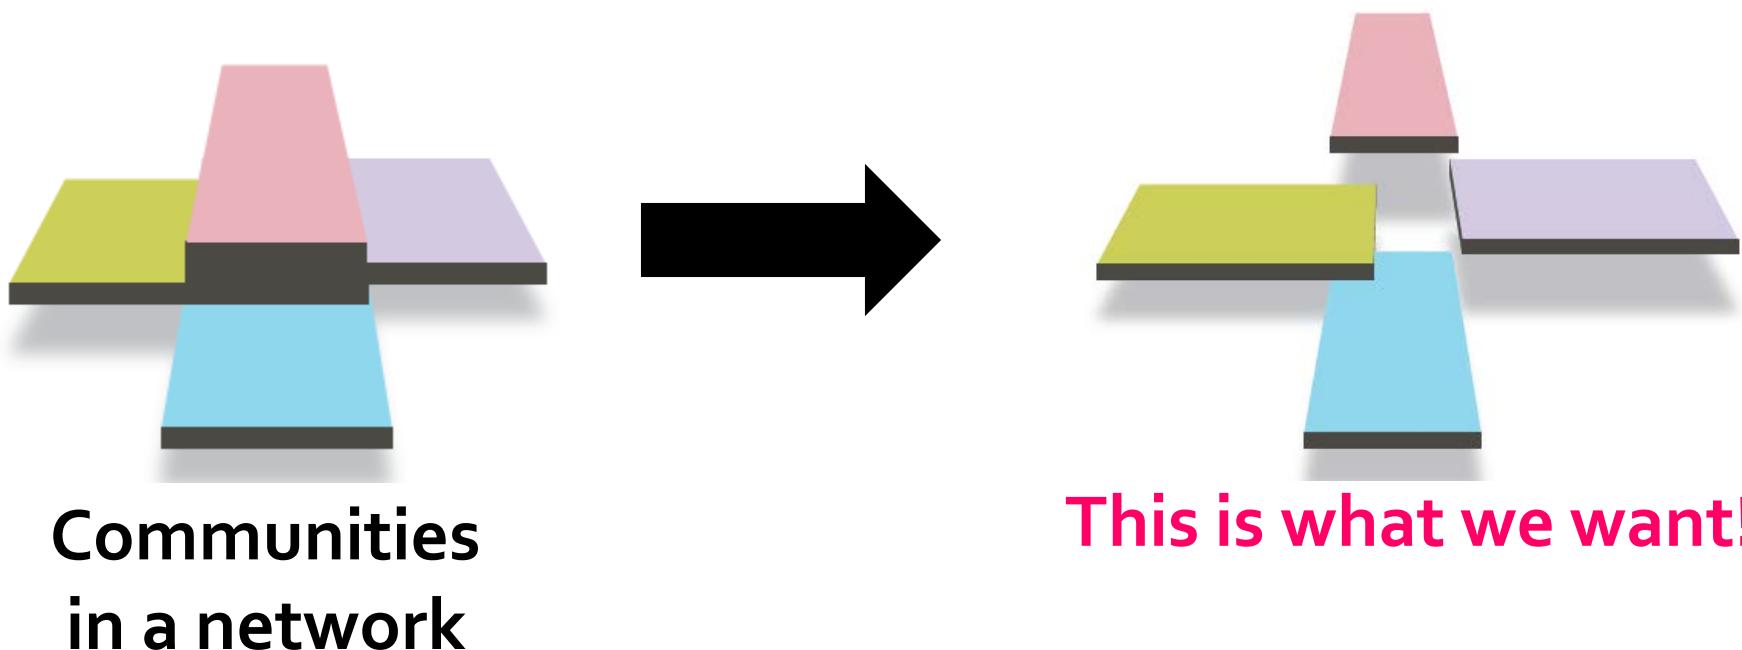
Communities as Tiles!

- What is the structure of community overlaps:
Edge density in the overlaps is higher!



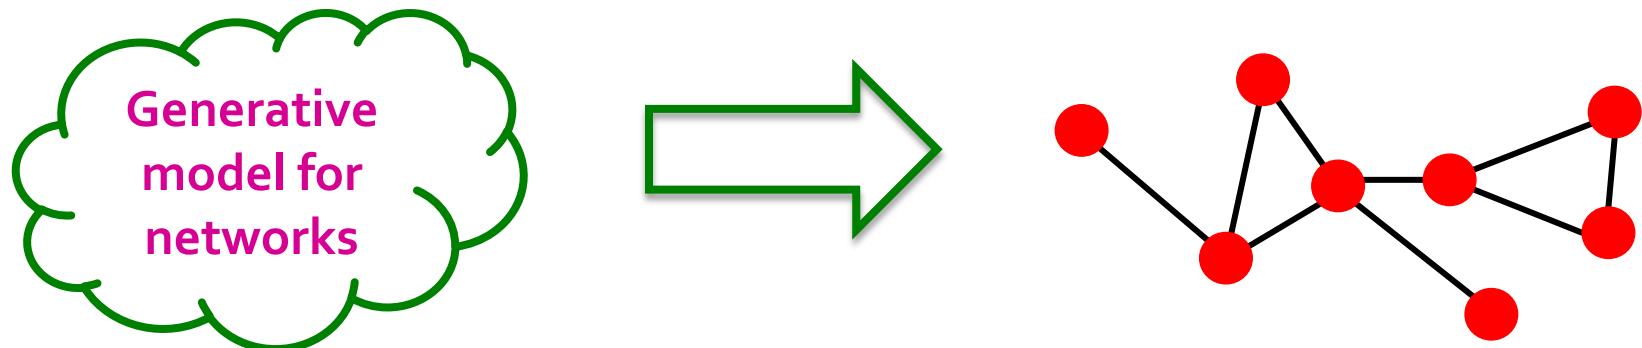
Communities as “tiles”

Recap so far...

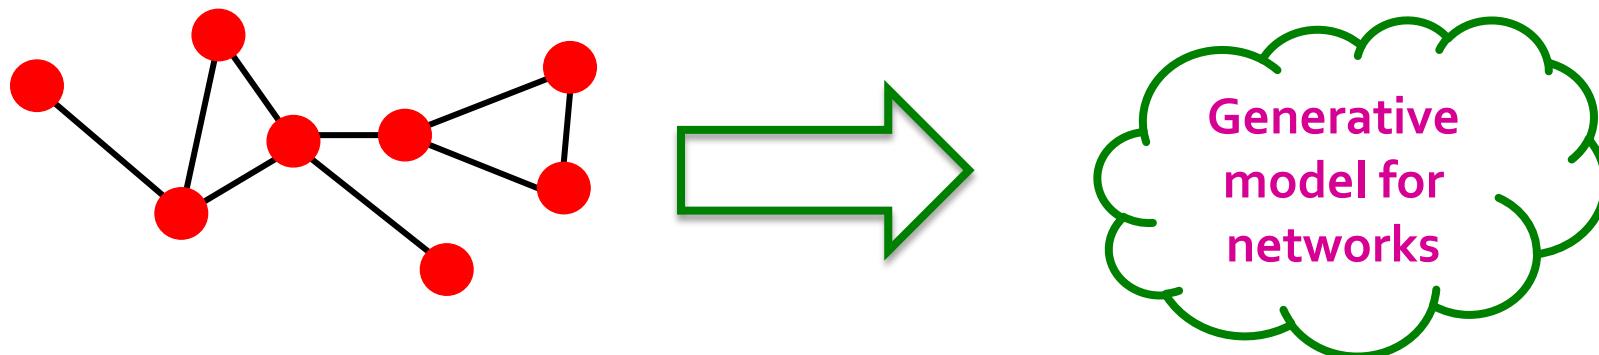


Plan of attack

- 1) Given a model, we generate the network:

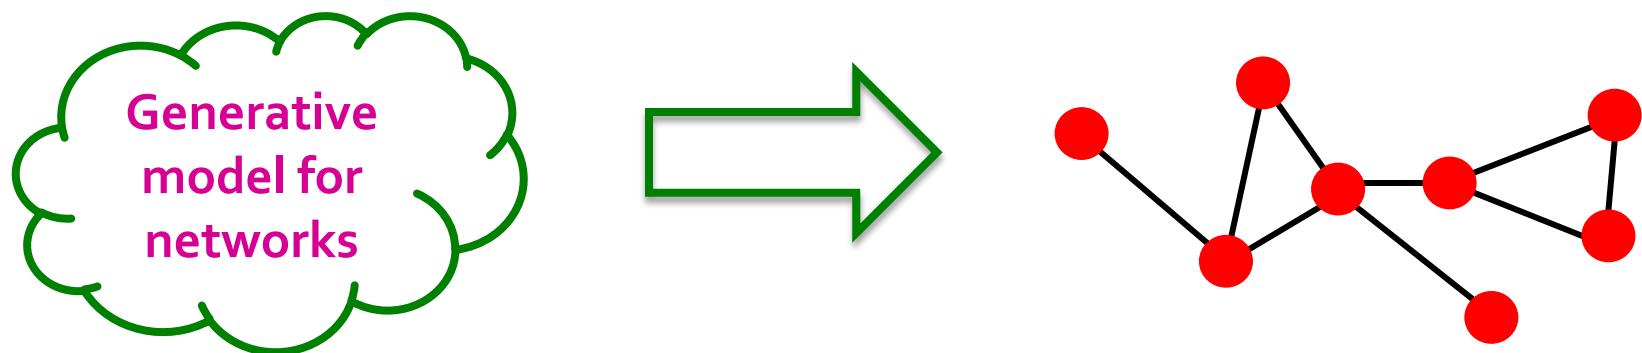


- 2) Given a network, find the “best” model



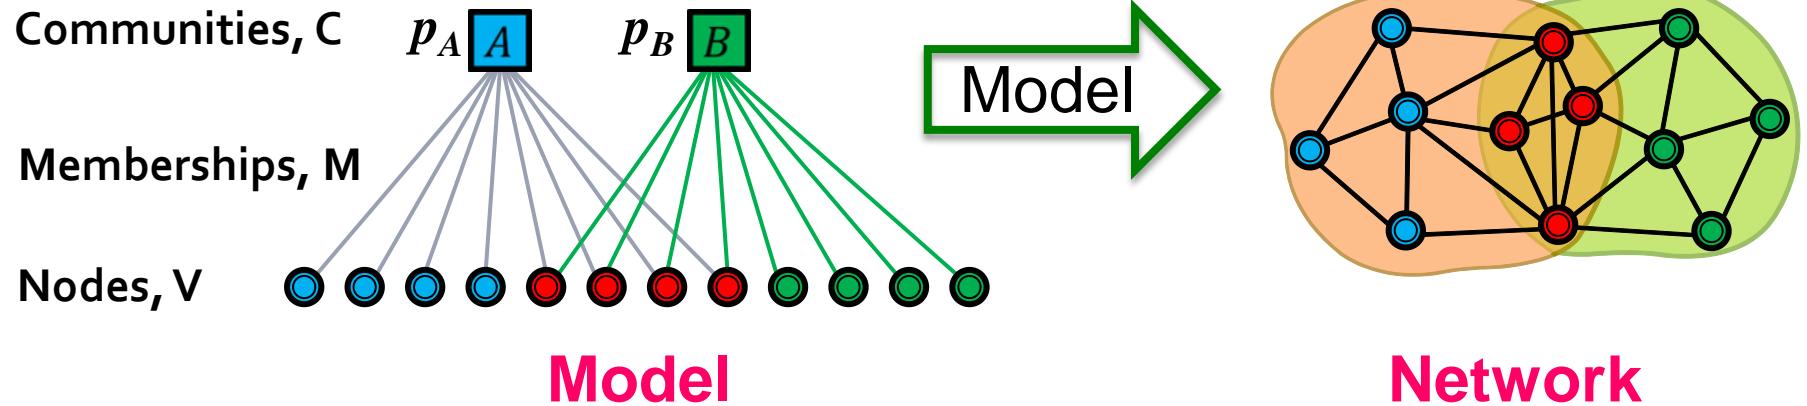
Model of networks

- Goal: Define a model that can generate networks
 - The model will have a set of “parameters” that we will later want to estimate (and detect communities)



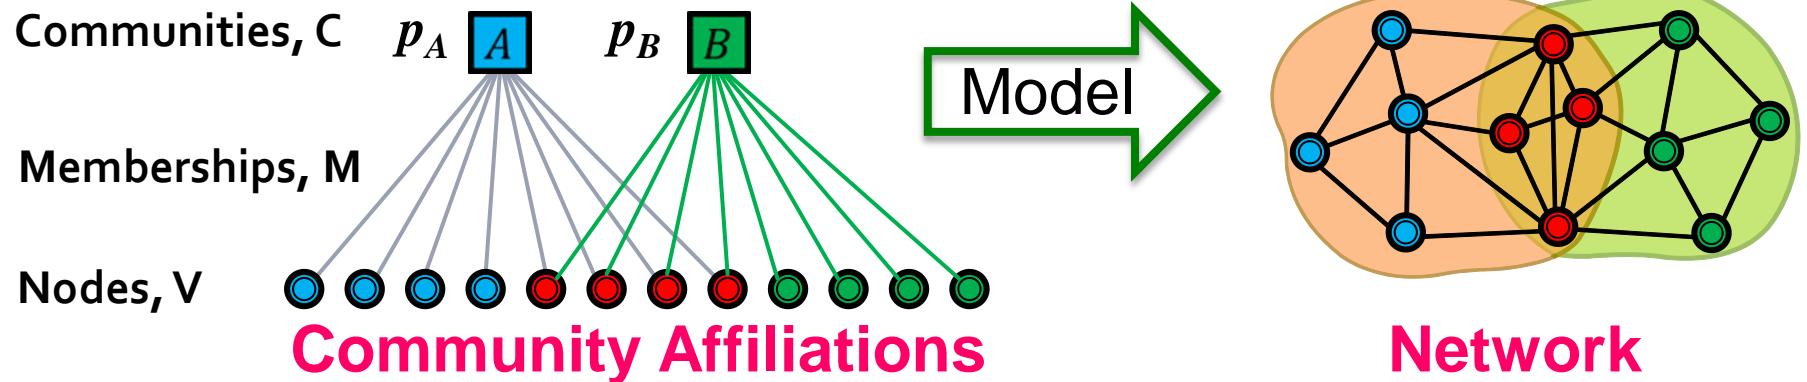
- Q: Given a set of nodes, how do communities “generate” edges of the network?

Community-Affiliation Graph



- **Generative model $B(V, C, M, \{p_c\})$ for graphs:**
 - Nodes V , Communities C , Memberships M
 - Each community c has a single probability p_c
 - Later we fit the model to networks to detect communities

AGM: Generative Process



- **AGM generates the links: For each**
 - For each pair of nodes in community A , we connect them with prob. p_A
 - **The overall edge probability is:**

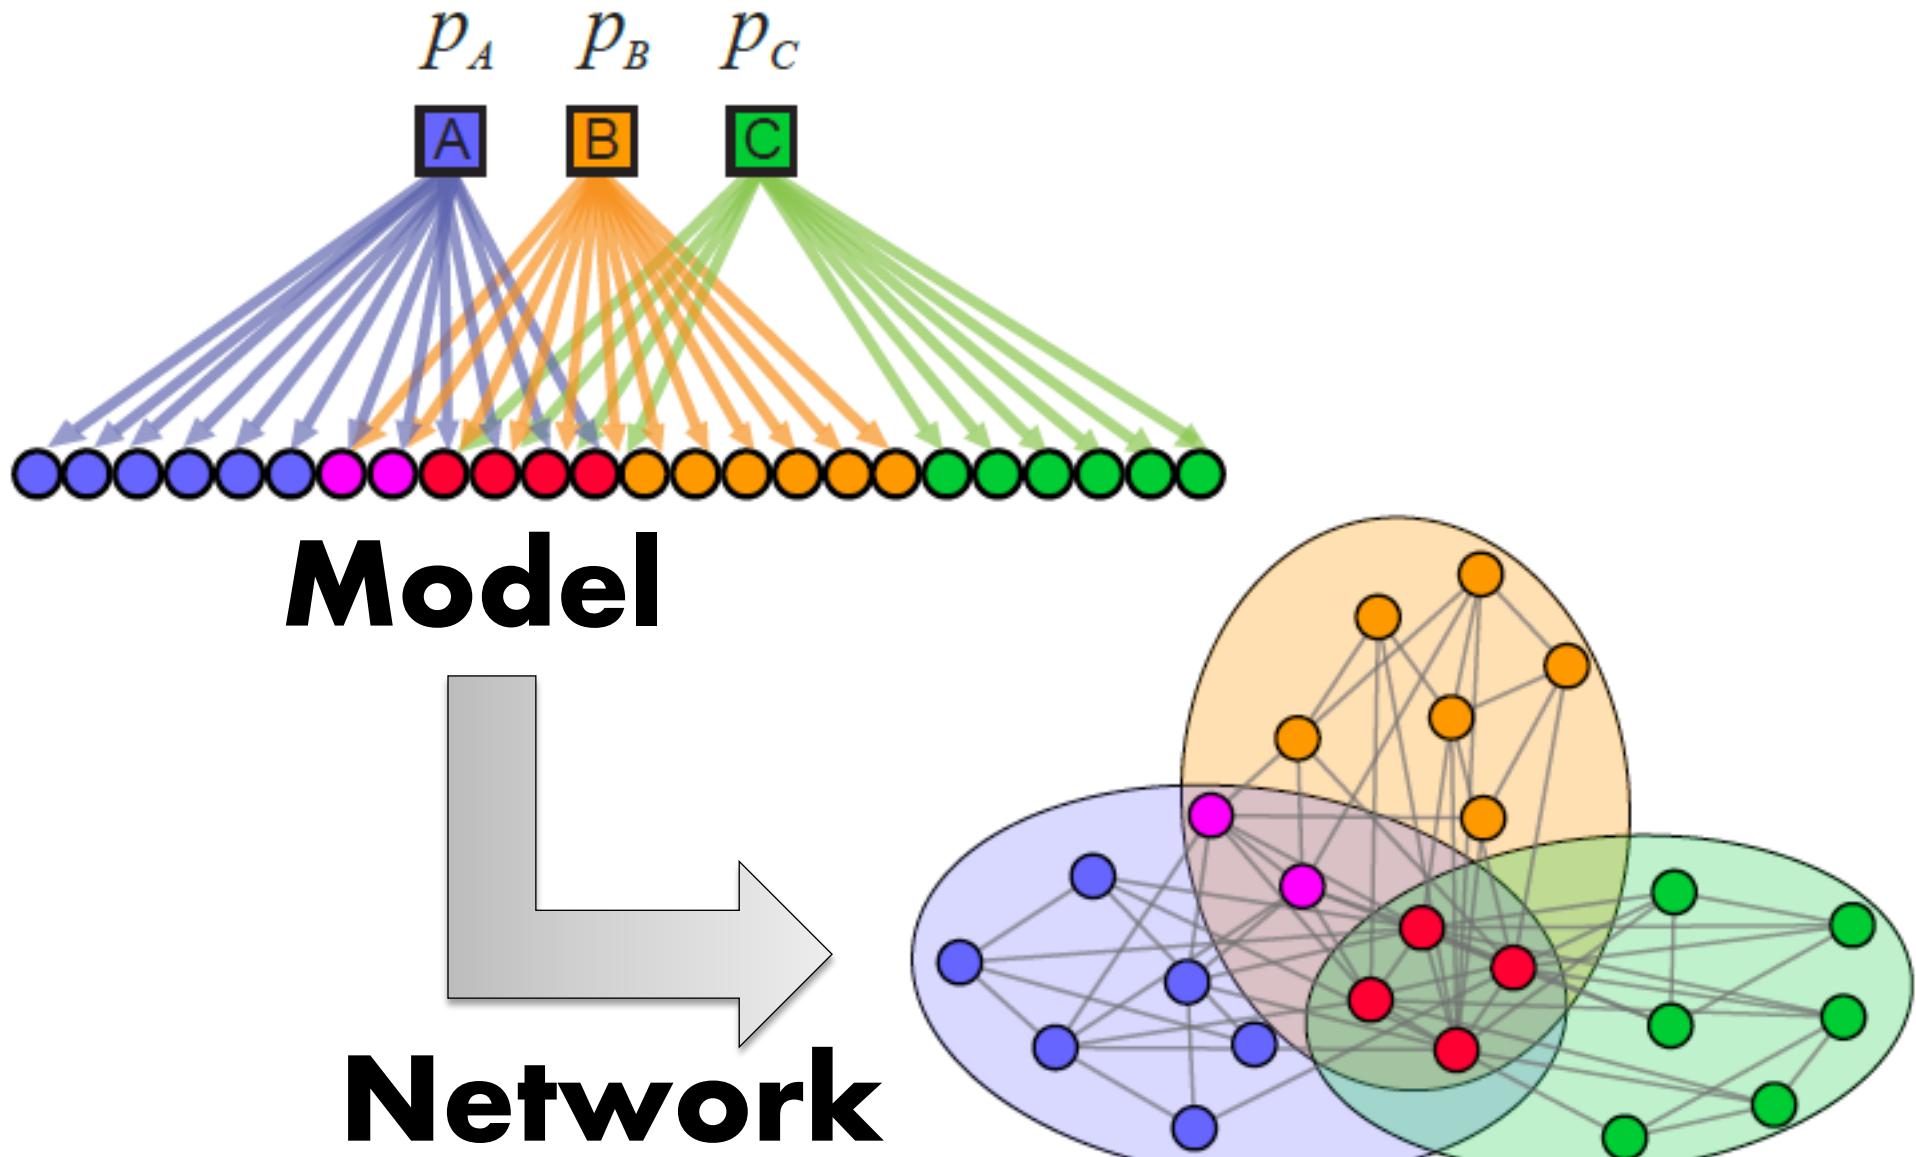
$$P(u, v) = 1 - \prod_{c \in M_u \cap M_v} (1 - p_c)$$

If u, v share no communities: $P(u, v) = \epsilon$

M_u ... set of communities
node u belongs to

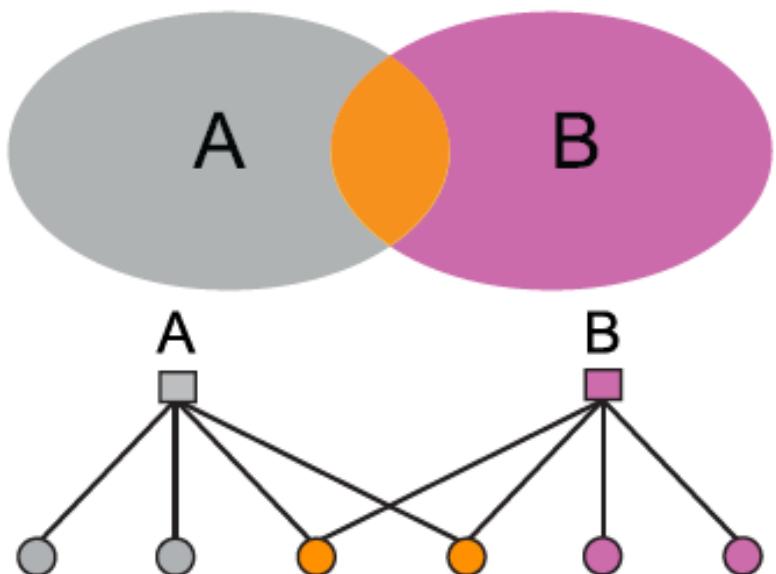
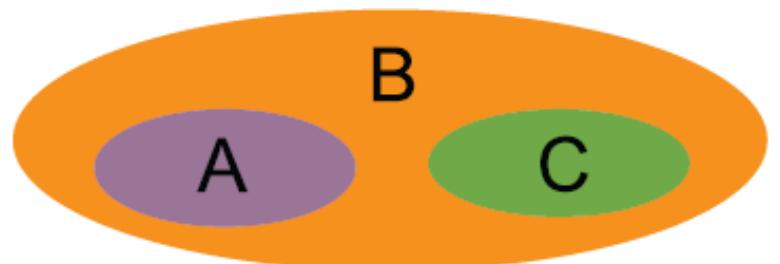
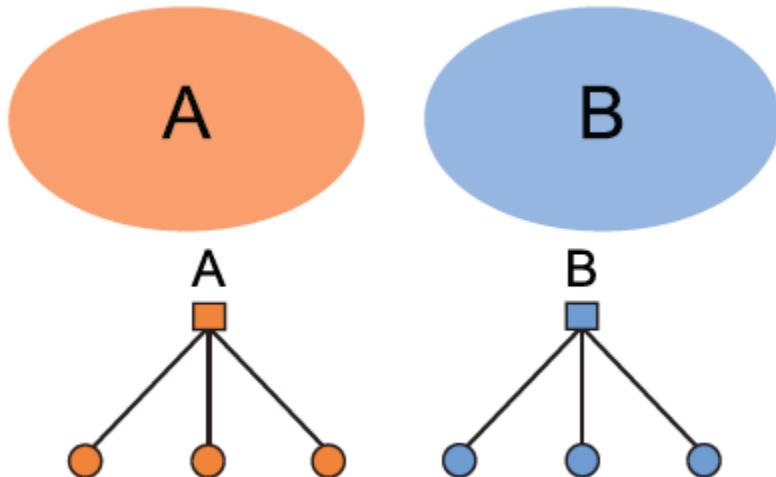
Think of this as an “OR” function: If at least 1 community says “YES” we create an edge

Recap: AGM networks



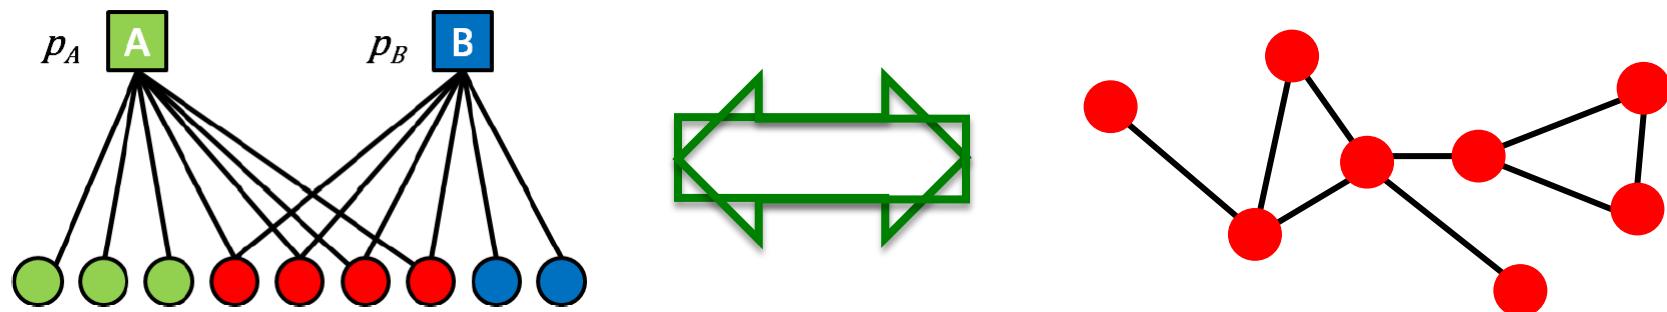
AGM: Flexibility

- AGM can express a variety of community structures:
Non-overlapping,
Overlapping, Nested



Detecting Communities

- Detecting communities with AGM:



Given a Graph, find the Model

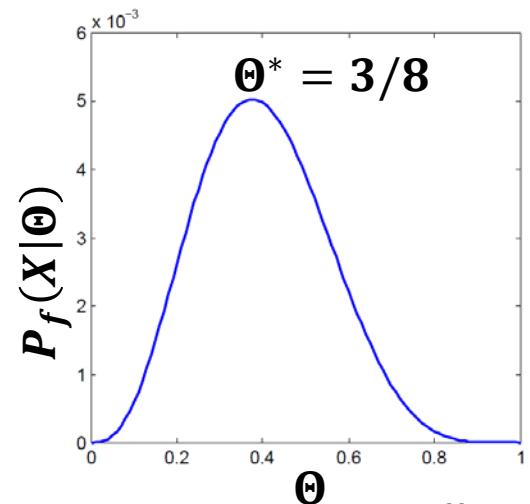
- 1) Affiliation graph M
- 2) Number of communities C
- 3) Parameters p_c

Maximum Likelihood Estimation

- **Maximum Likelihood Principle (MLE):**
 - **Given:** Data X
 - **Assumption:** Data is generated by some model $f(\Theta)$
 - f ... model
 - Θ ... model parameters
 - Want to estimate $P_f(X|\Theta)$:
 - The probability that our model f (with parameters Θ) generated the data
 - **Now let's find the most likely model that could have generated the data:** $\arg \max_{\Theta} P_f(X|\Theta)$

Example: MLE

- Imagine we are given a set of coin flips
- Task: Figure the bias of a coin!
 - Data: Sequence of coin flips: $X = [1, 0, 0, 0, 1, 0, 0, 1]$
 - Model: $f(\Theta)$ = return 1 with prob. Θ , else return 0
 - What is $P_f(X|\Theta)$? Assuming coin flips are independent
 - So, $P_f(X|\Theta) = P_f(0|\Theta) * P_f(1|\Theta) * P_f(1|\Theta) \dots * P_f(0|\Theta)$
 - What is $P_f(1|\Theta)$? Simple, $P_f(0|\Theta) = \Theta$
 - Then, $P_f(X|\Theta) = \Theta^3(1 - \Theta)^5$
 - For example:
 - $P_f(X|\Theta = 0.5) = 0.003906$
 - $P_f(X|\Theta = \frac{3}{8}) = 0.005029$
 - What did we learn? Our data was most likely generated by coin with bias $\Theta = 3/8$



MLE for Graphs

- How do we do MLE for graphs?
 - Model generates a **probabilistic adjacency matrix**
 - We then flip all the entries of the probabilistic matrix to obtain the **binary adjacency matrix**

For every pair
of nodes u, v
AGM gives the
prob. p_{uv} of
them being
linked

0.25	0.10	0.10	0.04
0.05	0.15	0.02	0.06
0.05	0.02	0.15	0.06
0.01	0.03	0.03	0.09

Flip
biased
coins

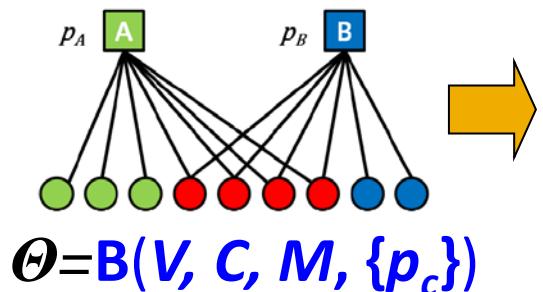
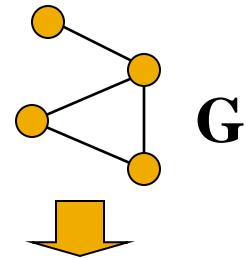
1	1	0	0
0	1	0	1
1	0	1	1
0	1	0	1

- The likelihood of AGM generating graph G:

$$P(G | \Theta) = \prod_{(u,v) \in G} P(u,v) \prod_{(u,v) \notin G} (1 - P(u,v))$$

Graphs: Likelihood $P(G|\Theta)$

- Given graph G and Θ we calculate likelihood that Θ generated G : $P(G|\Theta)$



0.25	0.10	0.10	0.04
0.05	0.15	0.02	0.06
0.05	0.02	0.15	0.06
0.01	0.03	0.03	0.09

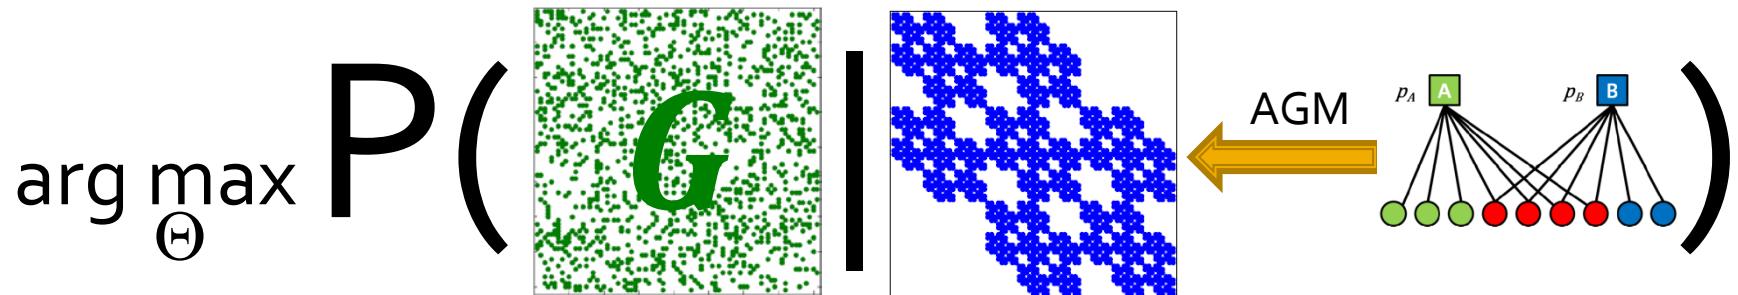
1	0	1	1
0	1	0	1
1	0	1	1
1	1	1	1

$$P(G|\Theta)$$

$$P(G | \Theta) = \prod_{(u,v) \in G} P(u,v) \prod_{(u,v) \notin G} (1 - P(u,v))$$

MLE for Graphs

- Our goal: Find $\Theta = B(V, C, M, \{p_C\})$ such that:

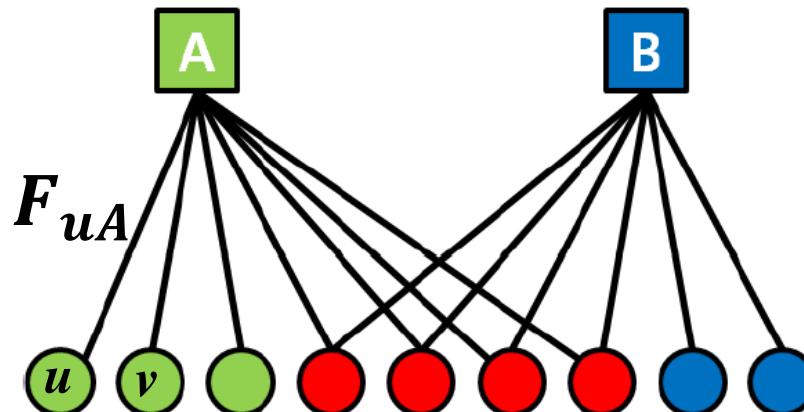


- How do we find $B(V, C, M, \{p_C\})$ that maximizes the likelihood?
 - No nice way, have to search over all affiliation graphs ☹ But, don't despair...



From AGM to BigCLAM

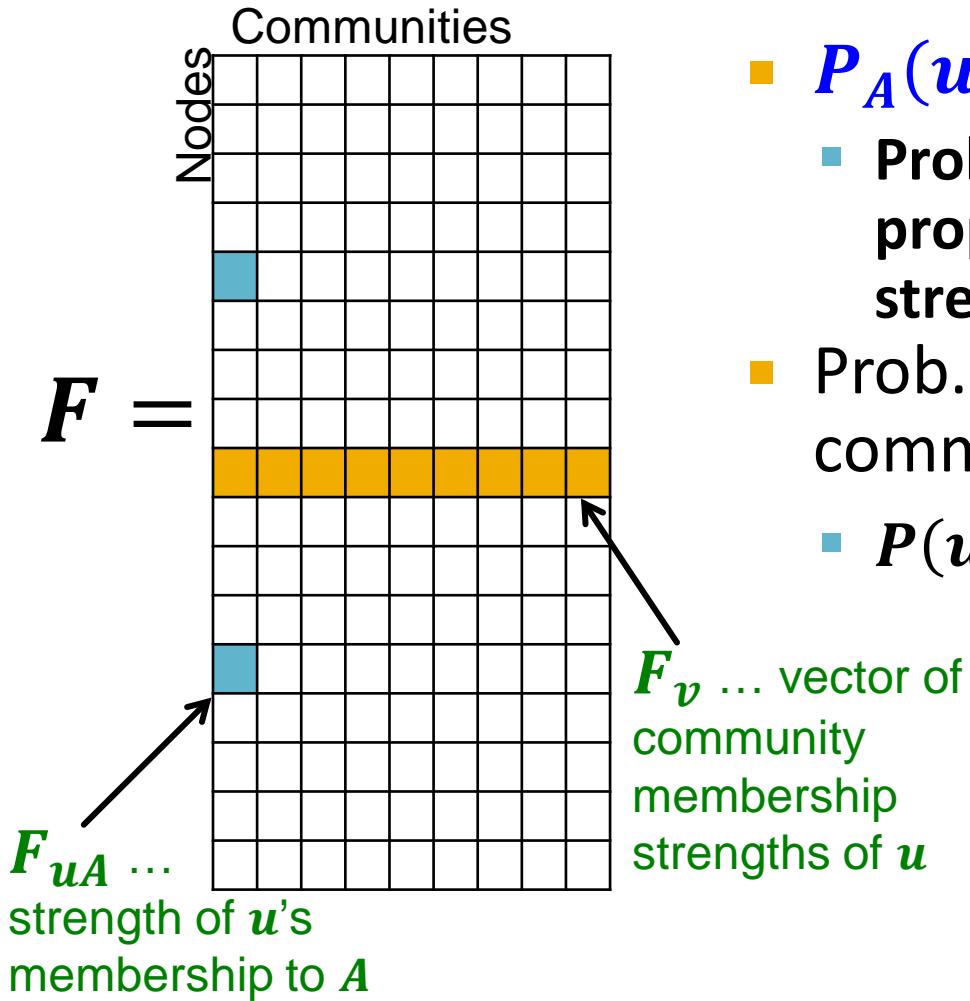
- Relaxation: Memberships have strengths



- F_{uA} : The membership strength of node u to community A ($F_{uA} = 0$: no membership)
- Each community A links nodes independently:
$$P_A(u, v) = 1 - \exp(-F_{uA} \cdot F_{vA})$$

From AGM to BigCLAM

■ Community membership strength matrix F



- $P_A(u, v) = 1 - \exp(-F_{uA} \cdot F_{vA})$
 - Probability of connection is proportional to the product of strengths
- Prob. that **at least one** common community C links the nodes:
 - $P(u, v) = 1 - \prod_C (1 - P_C(u, v))$

From AGM to BigCLAM

- Community A links nodes u, v independently:

$$P_A(u, v) = 1 - \exp(-F_{uA} \cdot F_{vA})$$

- Then prob. at least one common C links them:

$$\begin{aligned} P(u, v) &= 1 - \prod_C (1 - P_C(u, v)) \\ &= 1 - \exp(-\sum_C F_{uC} \cdot F_{vC}) \\ &= 1 - \exp(-F_u \cdot F_v^T) \end{aligned}$$

- For example:

$$F_u: \begin{array}{|c|c|c|c|} \hline 0 & 1.2 & 0 & 0.2 \\ \hline \end{array}$$

$$F_v: \begin{array}{|c|c|c|c|} \hline 0.5 & 0 & 0 & 0.8 \\ \hline \end{array}$$

$$F_w: \begin{array}{|c|c|c|c|} \hline 0 & 1.8 & 1 & 0 \\ \hline \end{array}$$

Node community
membership strengths

Then: $F_u \cdot F_v^T = 0.16$

And: $P(u, v) = 1 - \exp(-0.16) = 0.14$

But: $P(u, w) = 0.88$

$P(v, w) = 0$

From AGM to BigCLAM

- **Another way to think about this is:**
 - Think of a weighted graph (but we don't see the weights)
 - u, v interact with strength $X_{uv}^A \sim Pois(\mathbf{F}_{uA} \cdot \mathbf{F}_{vA})$
 - Total interaction strength is the sum: $X_{uv} = \sum_c X_{uv}^c$
 - Then: $X_{uv} \sim Pois(\sum_c \mathbf{F}_{uc} \cdot \mathbf{F}_{vc})$
 - **We observe an edge if there is non-zero interaction:**
$$P(X_{uv} > 0) = 1 - P(X_{uv} = 0) = \\ = 1 - \exp(-\mathbf{F}_u \cdot \mathbf{F}_v^T)$$

Poisson PMF:

$$P(X = x) = \frac{\lambda^k}{k!} e^{-\lambda}$$

BigCLAM

- **Task:** Given a network, estimate F
 - Find F that maximizes the log-likelihood
 - Many times we take the logarithm of the likelihood and call it log-likelihood: $l(F) = \log P(G|F)$

$$l(F) = \sum_{(u,v) \in E} \log(1 - \exp(-F_u F_v^T)) - \sum_{(u,v) \notin E} F_u F_v^T$$

- Objective function is biconvex
- **Non-negative matrix factorization:**
 - Update F_{uc} for node u while fixing the memberships of all other nodes

BigCLAM

- **Task:** Given a network, estimate F
 - Find F that maximizes the log-likelihood
 - Many times we take the logarithm of the likelihood and call it log-likelihood: $l(F) = \log P(G|F)$

$$l(F) = \sum_{(u,v) \in E} \log(1 - \exp(-F_u F_v^T)) - \sum_{(u,v) \notin E} F_u F_v^T$$

- **Connection:** Non-negative matrix factorization
 - We want to “factor” the adjacency matrix G

$$1 - \exp(-\begin{matrix} F \\ \cdot \\ F^T \end{matrix}) = \begin{matrix} \text{Matrix of edge probs.} \\ \rightarrow \\ \text{Graph } G \end{matrix}$$

How to find MLE

$$l(F) = \sum_{(u,v) \in E} \log(1 - \exp(-F_u F_v^T)) - \sum_{(u,v) \notin E} F_u F_v^T$$

- **Non-negative matrix factorization:**
 - Use a gradient based approach!
 - But updating the whole F takes too long
 - Computing the gradient takes quadratic time!
 - Our network are far too big to allow for that:
 - Network with 1M nodes, can have 10^{12} different edges
 - Instead, update F in small “units”:
 - Update F_{uC} for node u while fixing the memberships of all other nodes

BigCLAM: V1.0

$$l(F_u) = \sum_{v \in \mathcal{N}(u)} \log(1 - \exp(-F_u F_v^T)) - \sum_{v \notin \mathcal{N}(u)} F_u F_v^T$$

- Compute gradient of a single row F_u of F :

$$\nabla l(F_u) = \sum_{v \in \mathcal{N}(u)} F_v \frac{\exp(-F_u F_v^T)}{1 - \exp(-F_u F_v^T)} - \sum_{v \notin \mathcal{N}(u)} F_v$$

- Coordinate gradient ascent:

- Iterate over the rows of F :

- Compute gradient $\nabla l(F_u)$ of row u (while keeping others fixed)
- Update the row F_u : $F_u \leftarrow F_u + \eta \nabla l(F_u)$
- Project F_u back to a non-negative vector: If $F_{uC} < 0$: $F_{uC} = 0$

- This is slow! Computing $\nabla l(F_u)$ takes linear time!

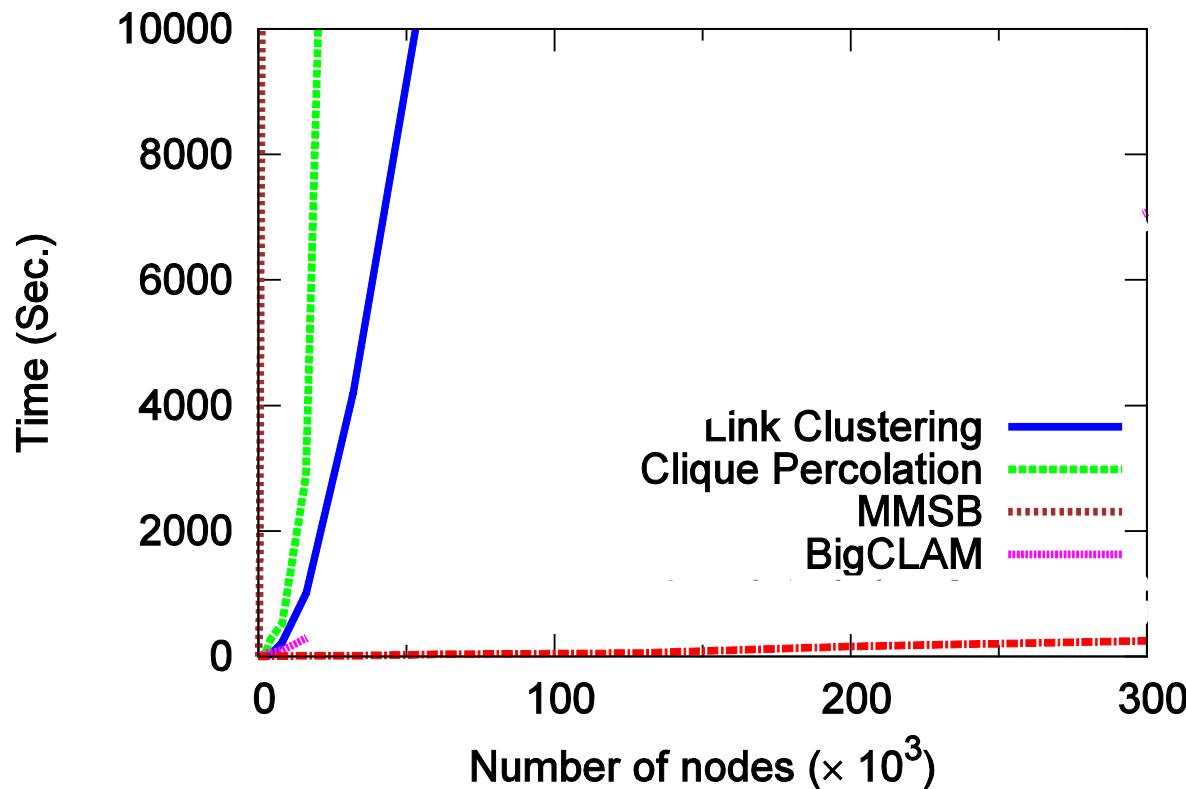
BigCLAM: V2.0

- However, we notice:

$$\sum_{v \notin \mathcal{N}(u)} F_v = \left(\sum_v F_v - F_u - \sum_{v \in \mathcal{N}(u)} F_v \right)$$

- We can cache $\sum_v F_v$
- So, computing $\sum_{v \notin \mathcal{N}(u)} F_v$ now takes **linear time** in the degree $\mathcal{N}(u)$ of u
 - In networks degree of a node is much smaller to the total number of nodes in the network, so this is a significant speedup!

BigClam: Scalability

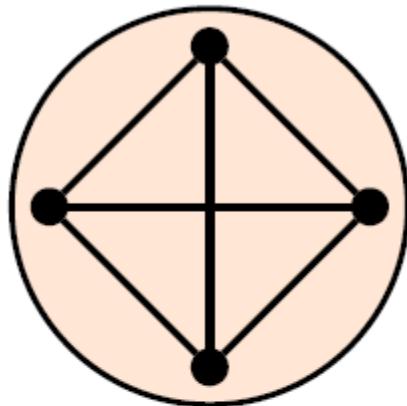


- **BigCLAM takes 5 minutes for 300k node nets**
 - Other methods take 10 days
- **Can process networks with 100M edges!**

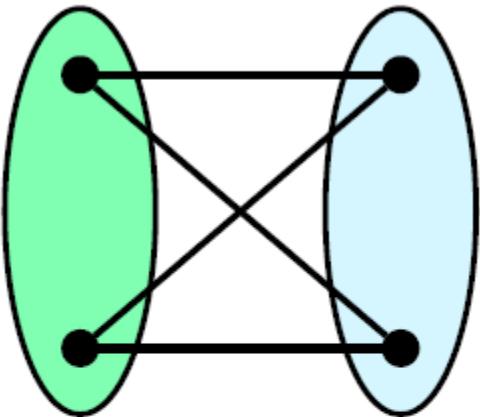
Extension: Beyond Clusters

Undirected

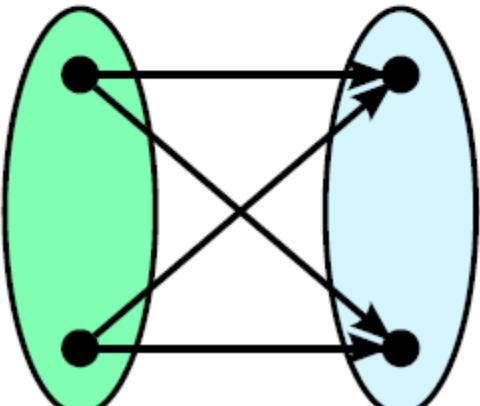
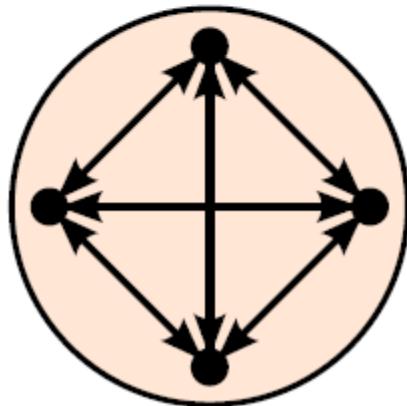
Cohesive



2-mode

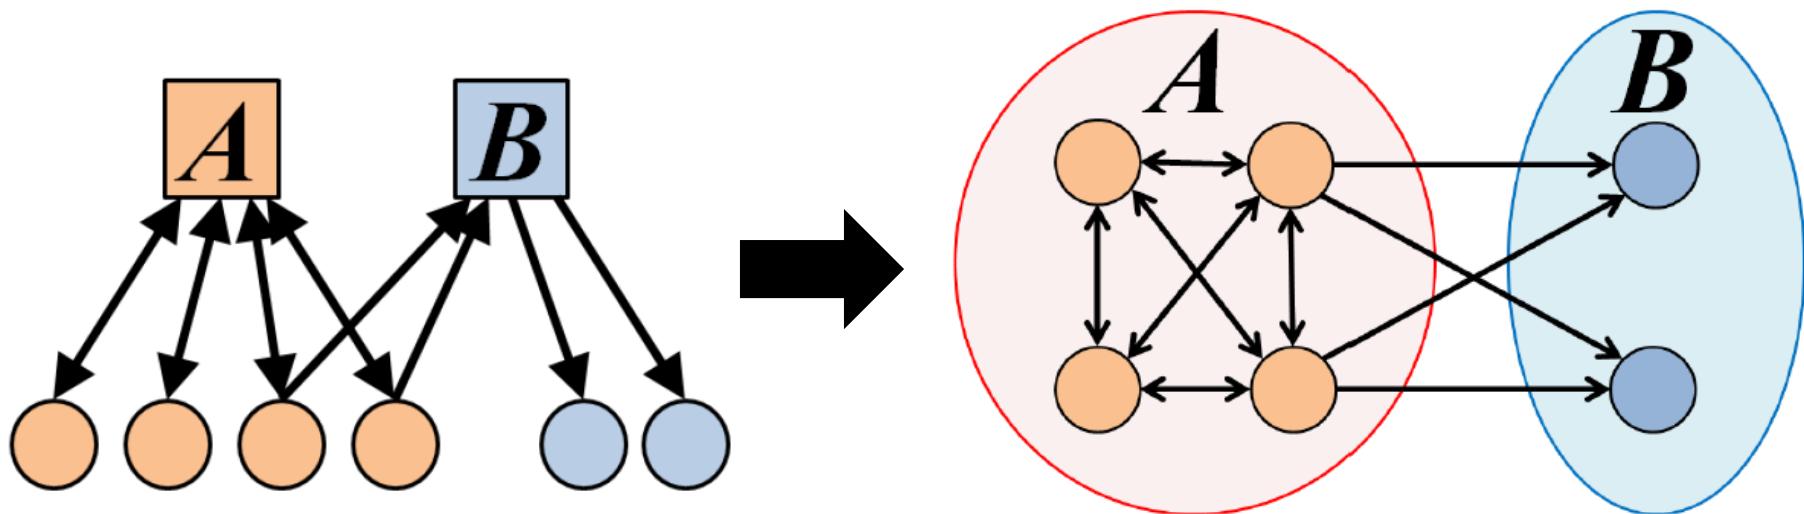


Directed

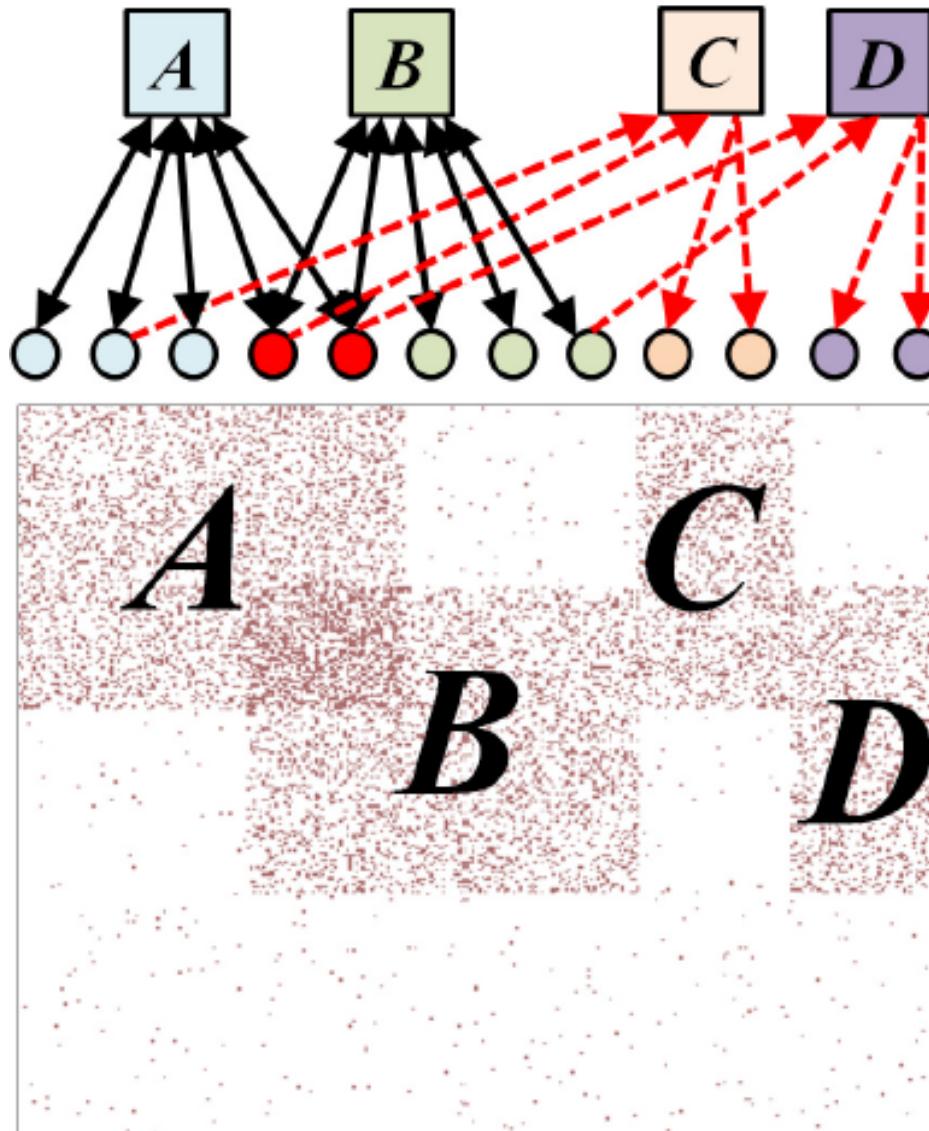


Extension: Directed AGM

- Extension: Make community membership edges directed:
 - Outgoing membership: Nodes “**sends**” edges
 - Incoming membership: Node “**receives**” edges



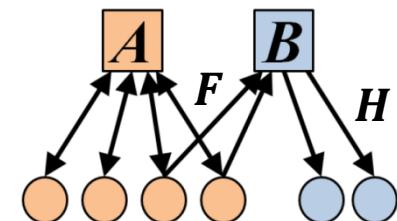
Example: Model and Network



Directed AGM

- Everything is almost the same except now we have 2 matrices: F and H

- F ... out-going community memberships
 - H ... in-coming community memberships

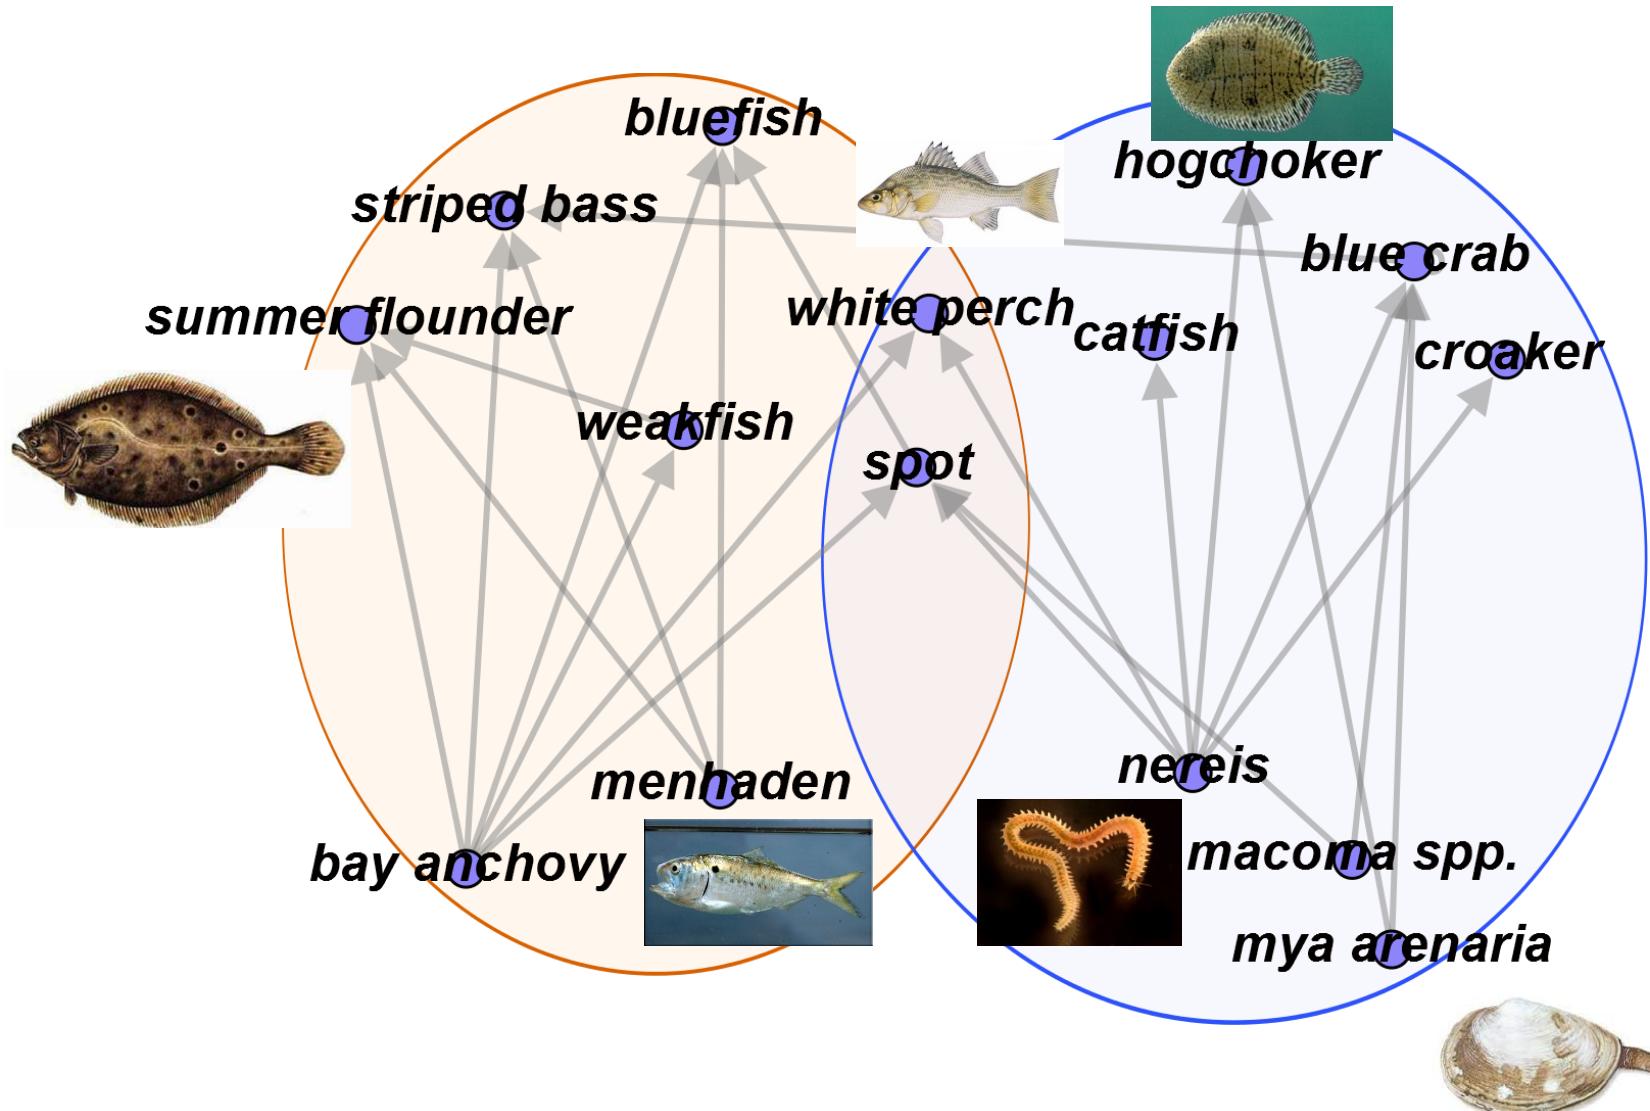


- Edge prob.: $P(u, v) = 1 - \exp(-F_u H_v^T)$
- Network log-likelihood:

$$l(F, H) = \sum_{(u, v) \in E} \log(1 - \exp(-F_u H_v^T)) - \sum_{(u, v) \notin E} F_u H_v^T$$

which we optimize the same way as before

Predator-prey Communities



More details at...

- [Overlapping Community Detection at Scale: A Nonnegative Matrix Factorization Approach](#) by J. Yang, J. Leskovec. *ACM International Conference on Web Search and Data Mining (WSDM)*, 2013.
- [Detecting Cohesive and 2-mode Communities in Directed and Undirected Networks](#) by J. Yang, J. McAuley, J. Leskovec. *ACM International Conference on Web Search and Data Mining (WSDM)*, 2014.
- [Community Detection in Networks with Node Attributes](#) by J. Yang, J. McAuley, J. Leskovec. *IEEE International Conference On Data Mining (ICDM)*, 2013.