

# Online Algorithms

## Performance-based Advertising

Mining of Massive Datasets  
Leskovec, Rajaraman, and Ullman  
Stanford University



# Online Algorithms

- **Classic model of algorithms**

- You get to see the entire input, then compute some function of it
- In this context, “offline algorithm”

- **Online Algorithms**

- You get to see the input one piece at a time, and need to make irrevocable decisions along the way
- **Similar to the data stream model**


# Online Algorithms

## Bipartite Matching

Mining of Massive Datasets  
Leskovec, Rajaraman, and Ullman  
Stanford University




# Example: Bipartite Matching




Nodes: Boys and Girls; Edges: Compatible Pairs  
**Goal: Match as many compatible pairs as possible**

# Example: Bipartite Matching



$M = \{(1,a), (2,b), (3,d)\}$  is a **matching**  
Cardinality of matching =  $|M| = 3$

# Example: Bipartite Matching



$M = \{(1,c), (2,b), (3,d), (4,a)\}$  is a  
**perfect matching**

**Perfect matching** ... all vertices of the graph are matched

**Maximum matching** ... a matching that contains the largest possible number of matches

# Matching Algorithm


- **Problem:** Find a maximum matching for a given bipartite graph
  - A perfect one if it exists
- There is a polynomial-time offline algorithm based on augmenting paths (Hopcroft & Karp 1973, see [http://en.wikipedia.org/wiki/Hopcroft-Karp\\_algorithm](http://en.wikipedia.org/wiki/Hopcroft-Karp_algorithm))
- **But what if we do not know the entire graph upfront?**

# Online Graph Matching Problem

- Initially, we are given the set boys
- In each round, one girl's choices are revealed
  - That is, girl's edges are revealed
- At that time, we have to decide to either:
  - Pair the girl with a boy
  - Do not pair the girl with any boy
- Example of application:  
Assigning tasks to servers



# Online Graph Matching: Example



(1,a)  
(2,b)  
(3,d)

# Greedy Algorithm

- Greedy algorithm for the online graph matching problem:
  - Pair the new girl with **any** eligible boy
    - If there is none, do not pair girl
- How good is the algorithm?

# Competitive Ratio

- For input  $I$ , suppose greedy produces matching  $M_{greedy}$  while an optimal matching is  $M_{opt}$


Competitive ratio =

$$\min_{\text{all possible inputs } I} (|M_{greedy}| / |M_{opt}|)$$

(what is greedy's worst performance over all possible inputs  $I$ )

# Analyzing the Greedy Algorithm


- Suppose  $M_{greedy} \neq M_{opt}$
- Consider the set  $G$  of girls matched in  $M_{opt}$  but not in  $M_{greedy}$
- (1)  $|M_{opt}| \leq |M_{greedy}| + |G|$
- Every boy  $B$  adjacent to girls in  $G$  is already matched in  $M_{greedy}$
- (2)  $|M_{greedy}| \geq |B|$



# Analyzing the Greedy Algorithm

## ■ So far:

- **G** matched in  $M_{opt}$  but not in  $M_{greedy}$
- Boys **B** adjacent to girls **G**
- (1)  $|M_{opt}| \leq |M_{greedy}| + |G|$
- (2)  $|M_{greedy}| \geq |B|$



- Optimal matches all the girls in **G** to boys in **B**
  - (3)  $|G| \leq |B|$
- Combining (2) and (3):
  - (4)  $|G| \leq |B| \leq |M_{greedy}|$


# Analyzing the Greedy Algorithm

- So we have:


- (1)  $|M_{\text{opt}}| \leq |M_{\text{greedy}}| + |G|$
- (4)  $|G| \leq |B| \leq |M_{\text{greedy}}|$

- Combining (1) and (4):

- $|M_{\text{opt}}| \leq |M_{\text{greedy}}| + |M_{\text{greedy}}|$
- $|M_{\text{opt}}| \leq 2|M_{\text{greedy}}|$
- $|M_{\text{greedy}}| / |M_{\text{opt}}| \geq 1/2$



# Worst-case Scenario



(1,a)  
(2,b)

# Online Algorithms

## Performance-based Advertising The AdWords Problem

Mining of Massive Datasets

Leskovec, Rajaraman, and Ullman

Stanford University





# History of Web Advertising

## ■ Banner ads (1995-2001)

- Initial form of web advertising
- Popular websites charged  $X\$$  for every 1,000 “impressions” of the ad
  - Called “CPM” rate  
(Cost per thousand impressions)
  - Modeled similar to TV, magazine ads
- From **untargeted** to **demographically targeted**
- **Low click-through rates**
  - Low ROI for advertisers



# Performance-based Advertising

- Introduced by Overture around 2000
  - Advertisers **bid on search keywords**
  - When someone searches for that keyword, the **highest bidder's ad is shown**
  - Advertiser is charged only if the ad is clicked on
- Similar model adopted by Google with some changes around 2002
  - Called **Adwords**

# Algorithmic Challenges

- **Performance-based advertising works!**
  - Multi-billion-dollar industry
- **What ads to show for a given query?**
  - (Today's lecture)
- **If I am an advertiser, which search terms should I bid on and how much should I bid?**
  - (Not focus of today's lecture)

# AdWords Problem

- A stream of queries arrives at the search engine:  $q_1, q_2, \dots$
- Several advertisers bid on each query
- When query  $q_i$  arrives, search engine must pick a subset of advertisers whose ads are shown
- **Goal:** Maximize search engine's revenues
- **Clearly we need an online algorithm!**

# Expected Revenue

Advertiser	Bid	CTR	Bid * CTR
A	\$1.00	1%	1 cent
B	\$0.75	2%	1.5 cents
C	\$0.50	2.5%	1.125 cents

Click through  
rate

Expected  
revenue

# The Adwords Innovation

Instead of sorting advertisers by bid, sort by expected revenue!

Advertiser	Bid	CTR	Bid * CTR
B	\$0.75	2%	1.5 cents
C	\$0.50	2.5%	1.125 cents
A	\$1.00	1%	1 cent

# Adwords Problem

## ■ Given:

- A set of bids by advertisers for search queries
- A click-through rate for each advertiser-query pair
- A budget for each advertiser (say for 1 day, month...)
- A limit on the number of ads to be displayed with each search query

## ■ Respond to each search query with a set of advertisers such that:

- The size of the set is no larger than the limit on the number of ads per query
- Each advertiser has bid on the search query
- Each advertiser has enough budget left to pay for the ad if it is clicked upon

# Limitations of Simple Algorithm

Instead of sorting advertisers by bid, sort by expected revenue!

Advertiser	Bid	CTR	Bid * CTR
B	\$0.75	2%	1.5 cents
C	\$0.50	2.5%	1.125 cents
A	\$1.00	1%	1 cent

- CTR of an ad is unknown
- Advertisers have limited budgets and bid on multiple ads (BALANCE algorithm)



# Estimating CTR

- Clickthrough rate (CTR) for a query-ad pair is measured historically
  - Averaged over a time period
- Some complications we won't cover in this lecture
  - CTR is position dependent
    - Ad #1 is clicked more than Ad #2
  - Explore v Exploit: Keep showing ads we already know the CTR of, or show new ads to estimate their CTR?

# Online Algorithms

## The BALANCE Algorithm

Mining of Massive Datasets

Leskovec, Rajaraman, and Ullman

Stanford University



# Adwords Problem

## ■ Given:

- A set of bids by advertisers for search queries
- A click-through rate for each advertiser-query pair
- A budget for each advertiser (say for 1 day, month...)
- A limit on the number of ads to be displayed with each search query

## ■ Respond to each search query with a set of advertisers such that:

- The size of the set is no larger than the limit on the number of ads per query
- Each advertiser has bid on the search query
- Each advertiser has enough budget left to pay for the ad if it is clicked upon

# Dealing with Limited Budgets

- **Our setting: Simplified environment**
  - There is **1** ad shown for each query
  - All advertisers have the same budget  **$B$**
  - All ads are equally likely to be clicked
  - Value of each ad is the same ( **$=1$** )
- **Simplest algorithm is greedy:**
  - For a query pick any advertiser who has bid **1** for that query
  - **Competitive ratio of greedy is  $1/2$**

# Bad Scenario for Greedy

- **Two advertisers A and B**
  - A bids on query  $x$ , B bids on  $x$  and  $y$
  - Both have budgets of \$4
- **Query stream:  $x x x x y y y y$** 
  - Worst case greedy choice:  $B B B B \_ \_ \_ \_$
  - Optimal:  $A A A A B B B B$
  - Competitive ratio =  $\frac{1}{2}$
- **This is the worst case!**
  - **Note:** Greedy algorithm is deterministic – it always resolves draws in the same way

# BALANCE Algorithm [MSVV]

- **BALANCE** Algorithm by Mehta, Saberi, Vazirani, and Vazirani
  - For each query, pick the advertiser with the largest unspent budget
  - Break ties arbitrarily (but in a deterministic way)

# Example: BALANCE

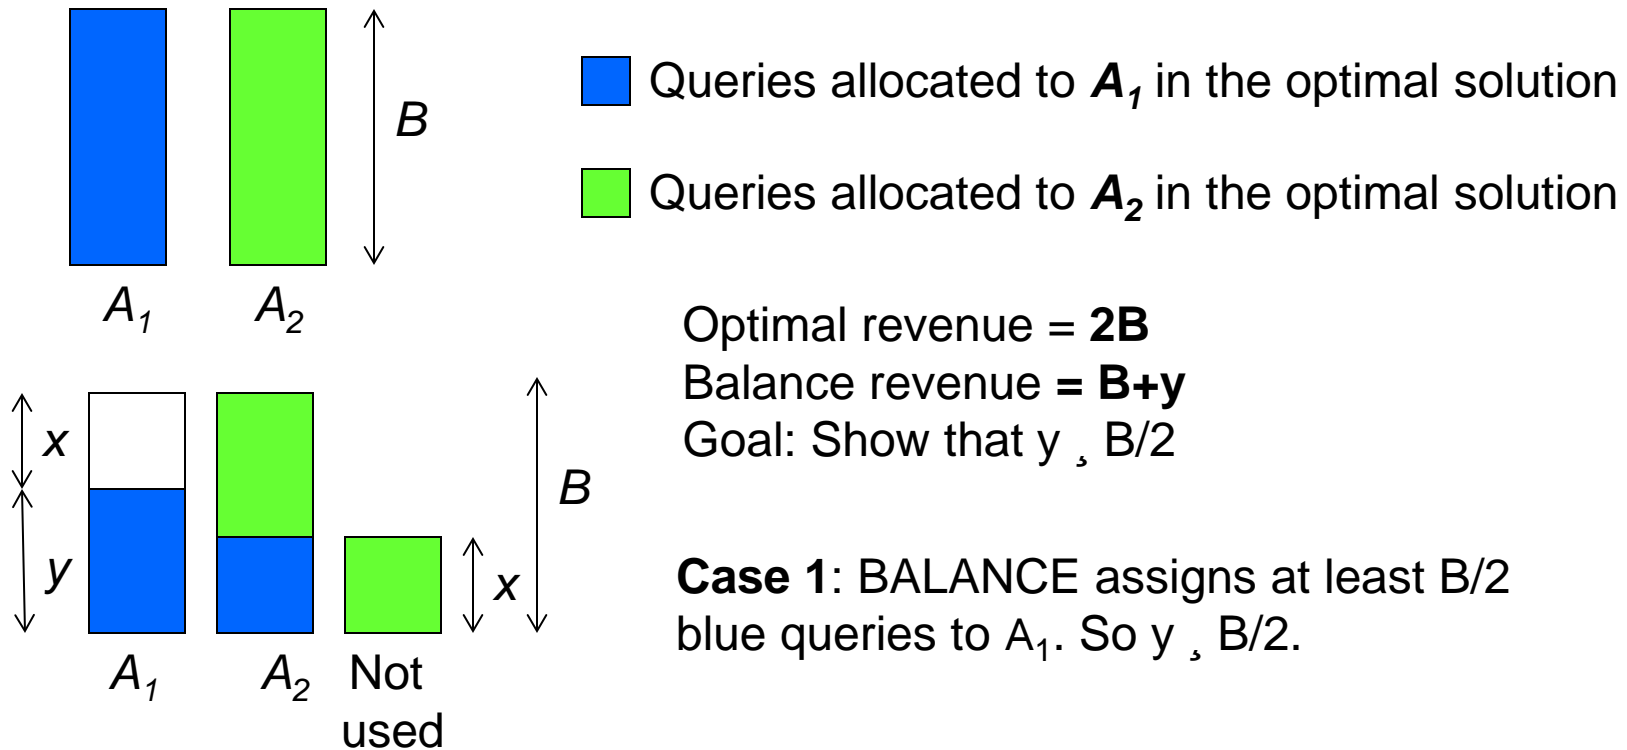
- **Two advertisers A and B**
  - A bids on query  $x$ , B bids on  $x$  and  $y$
  - Both have budgets of \$4
- **Query stream:**  $x\ x\ x\ x\ y\ y\ y\ y$
- **BALANCE choice:**  $A\ B\ A\ B\ B\ B\ \_\ \_$ 
  - Optimal:  $A\ A\ A\ A\ B\ B\ B\ B$
- **Competitive ratio =  $\frac{3}{4}$** 
  - For BALANCE with 2 advertisers

# Analyzing 2-advertiser BALANCE

- **Consider simple case**
  - 2 advertisers,  $A_1$  and  $A_2$ , each with budget  $B$  ( $\geq 1$ )
  - Optimal solution exhausts both advertisers' budgets
- **BALANCE must exhaust at least one advertiser's budget:**
  - If not, we can allocate more queries
  - Assume BALANCE exhausts  $A_2$ 's budget



# Analyzing Balance



**Case 2:** BALANCE assigns more than  $B/2$  blue queries to  $A_2$ .

Consider the last blue query assigned to  $A_2$ .

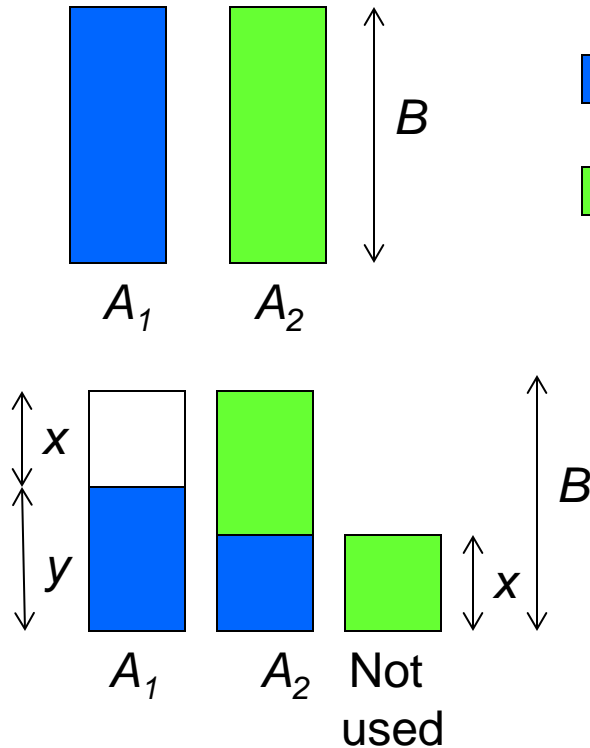
At that time,  $A_2$ 's unspent budget must have been at least as big as  $A_1$ 's.

That means at least as many queries have been assigned to  $A_1$  as to  $A_2$ .

At this point, we have already assigned at least  $B/2$  queries to  $A_2$ .

So  $y \geq B/2$ .

# Analyzing BALANCE



■ Queries allocated to  $A_1$  in the optimal solution

■ Queries allocated to  $A_2$  in the optimal solution

Optimal revenue  $OPT = 2B$

Balance revenue  $BAL = B + y$

We have shown that  $y \leq B/2$

$BAL \leq B + B/2 = 3B/2$

$BAL/OPT \leq 3/4$

# BALANCE: General Result

- In the general case, worst competitive ratio of BALANCE is  $1 - 1/e = \text{approx. } 0.63$ 
  - Interestingly, no online algorithm has a better competitive ratio!
- Let's see the worst case example that gives this ratio


# Worst case for BALANCE

- **$N$  advertisers:**  $A_1, A_2, \dots, A_N$ 
  - Each with budget  $B > N$
- **Queries:**
  - $N \cdot B$  queries appear in  $N$  rounds of  $B$  queries each
- **Bidding:**
  - Round 1 queries: bidders  $A_1, A_2, \dots, A_N$
  - Round 2 queries: bidders  $A_2, A_3, \dots, A_N$
  - Round  $i$  queries: bidders  $A_i, \dots, A_N$
- **Optimum allocation:**

Allocate round  $i$  queries to  $A_i$

  - Optimum revenue  $N \cdot B$

# BALANCE Allocation





After  $k$  rounds, the allocation to advertiser  $k$  is:

$$S_k = \sum_{i=1}^k B/(N-i+1)$$

**If we find the smallest  $k$  such that  $S_k \geq B$ , then after  $k$  rounds we cannot allocate any queries to any advertiser**


# BALANCE: Analysis



# BALANCE: Analysis

- **Fact:** for large  $n$ 
  - Result due to Euler

$1/1 \quad 1/2 \quad 1/3 \quad \dots \quad 1/(N-(k-1)) \quad \dots \quad 1/(N-1) \quad 1/N$



$$\ln(N-k) = \ln(N) - 1$$

$$\ln(N/(N-k)) = 1$$

$$N/(N-k) = e$$

$$k = N(1 - 1/e)$$

# BALANCE: Analysis

- So after the first  $k=N(1-1/e)$  rounds, we cannot allocate a query to any advertiser
- Revenue =  $B \cdot N (1-1/e)$
- Competitive ratio =  $1-1/e$



# General Version of the Problem

- So far: all bids = 1, all budgets equal (=B)
- In a general setting **BALANCE** can be terrible
  - Consider query  $\mathbf{q}$ , two advertisers  $\mathbf{A}_1$  and  $\mathbf{A}_2$
  - $\mathbf{A}_1$ : *bid* = 1, *budget* = 110
  - $\mathbf{A}_2$ : *bid* = 10, *budget* = 100
  - Suppose we see 10 instances of  $\mathbf{q}$
  - BALANCE always selects  $\mathbf{A}_1$  and earns 10
  - Optimal earns 100

# Generalized BALANCE

- Consider query  $q$ , bidder  $i$ 
  - Bid =  $x_i$
  - Budget =  $b_i$
  - Amount spent so far =  $m_i$
  - Fraction of budget left over  $f_i = 1 - m_i/b_i$
  - Define  $\psi_i(q) = x_i(1 - e^{-f_i})$
- Allocate query  $q$  to bidder  $i$  with largest value of  $\psi_i(q)$
- Same competitive ratio  $(1 - 1/e)$