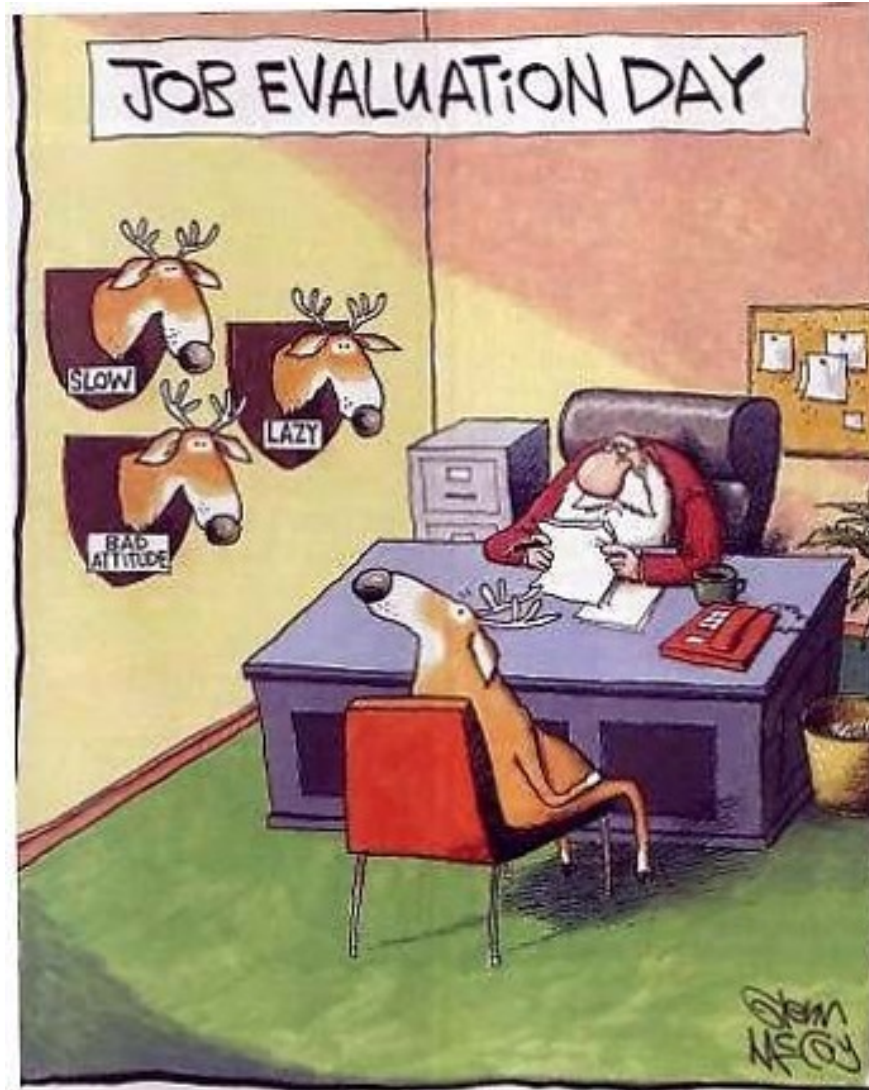


# Mining Massive Datasets: Review

CS246: Mining Massive Datasets  
Jure Leskovec, Stanford University  
<http://cs246.stanford.edu>



# Please fill out Course Evaluations!

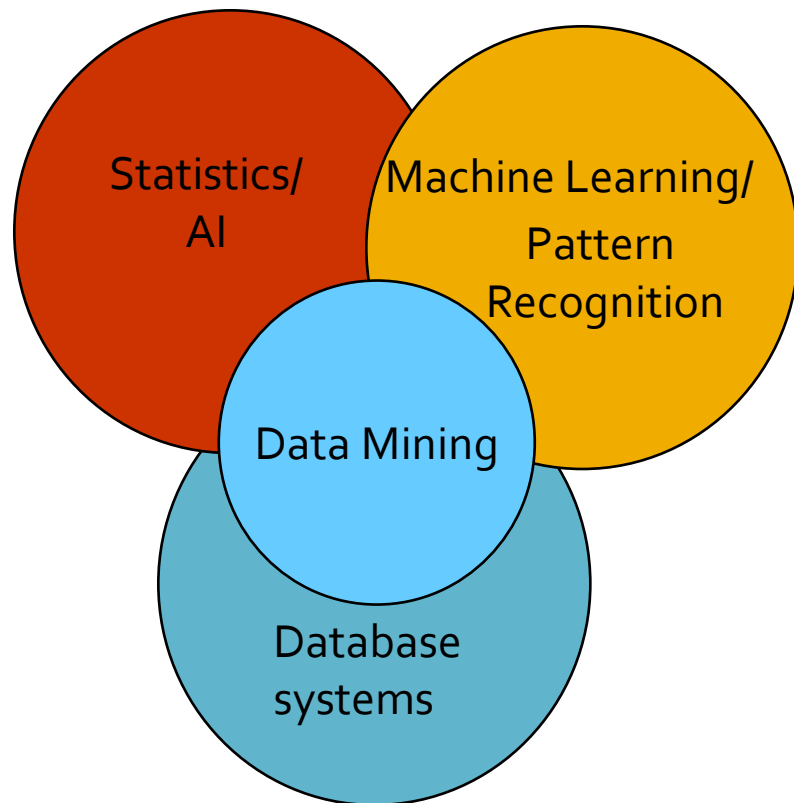


# Data Mining

- **Models and tools for discovering patterns and answering queries that are:**
  - **Valid:** Hold on new data with some certainty
  - **Useful:** Should be possible to act on the item
  - **Unexpected:** Non-obvious to the system
  - **Understandable:** Humans should be able to interpret the pattern

# Mining Massive Datasets

- Overlaps with machine learning, statistics, artificial intelligence, databases, but more stress on
  - **Scalability** of number of features and instances
  - **Algorithms** and **architectures**
  - Automation for handling **large data**



# What We Have Covered

- Apriori
- MapReduce
- Association rules
- Frequent itemsets
- PCY
- Recommender systems
- PageRank
- TrustRank
- HITS
- SVM
- Decision Trees
- Perceptron
- Web Advertising
- DGIM
- Bandits
- BFR
- Regret
- LSH
- MinHash
- SVD
- Clustering
- Matrix factorization
- CUR
- Bloom filters
- Flajolet-Martin
- CURE
- Submodularity
- SGD
- Collaborative Filtering
- SimRank
- Random hyperplanes
- Trawling
- AND-OR constructions
- k-means

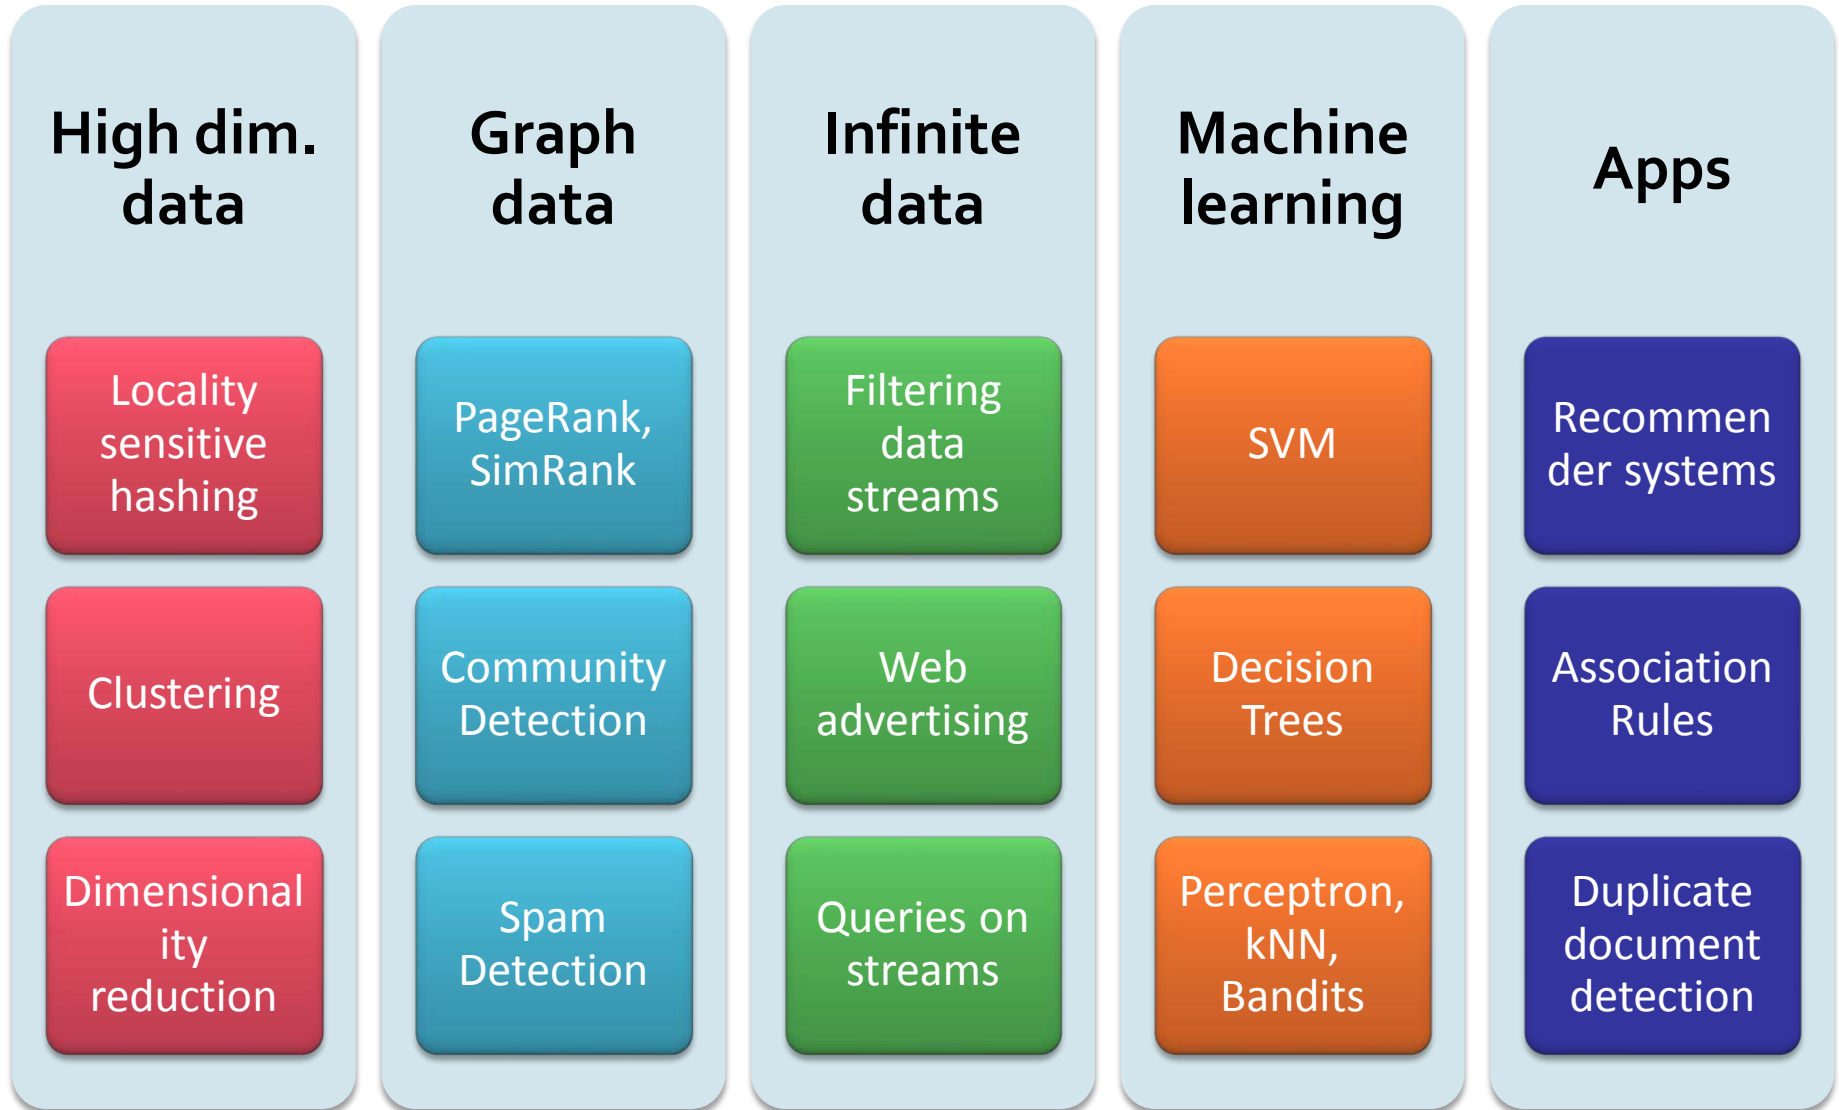
# How It All Fits Together

- **Based on different types of data:**
  - Data is **high dimensional**
  - Data is a **graph**
  - Data is **never-ending**
  - Data is **labeled**
- **Based on different models of computation:**
  - **MapReduce**
  - **Streams**
  - **Batch (offline) vs. Active (online) algorithms**
  - **Single machine in-memory**

# How It All Fits Together

- **Based on different applications:**
  - Recommender systems
  - Market basket analysis
  - Link analysis, spam detection
  - Duplicate detection and similarity search
  - Web advertising
- **Based on different “tools”:**
  - Linear algebra: SVD, Matrix factorization
  - Optimization: Stochastic gradient descent
  - Dynamic programming: Frequent itemsets
  - Hashing: LSH, Bloom filters,

# How It All Fits Together





# How it all fits together?

## Data is High-dimensional:

- Locality Sensitive Hashing
- Dimensionality reduction
- Clustering

## Data is a graph:

- Link Analysis: PageRank, TrustRank, Hubs & Authorities

## Data is Labeled (Machine Learning):

- kNN, Perceptron, SVM, Decision Trees

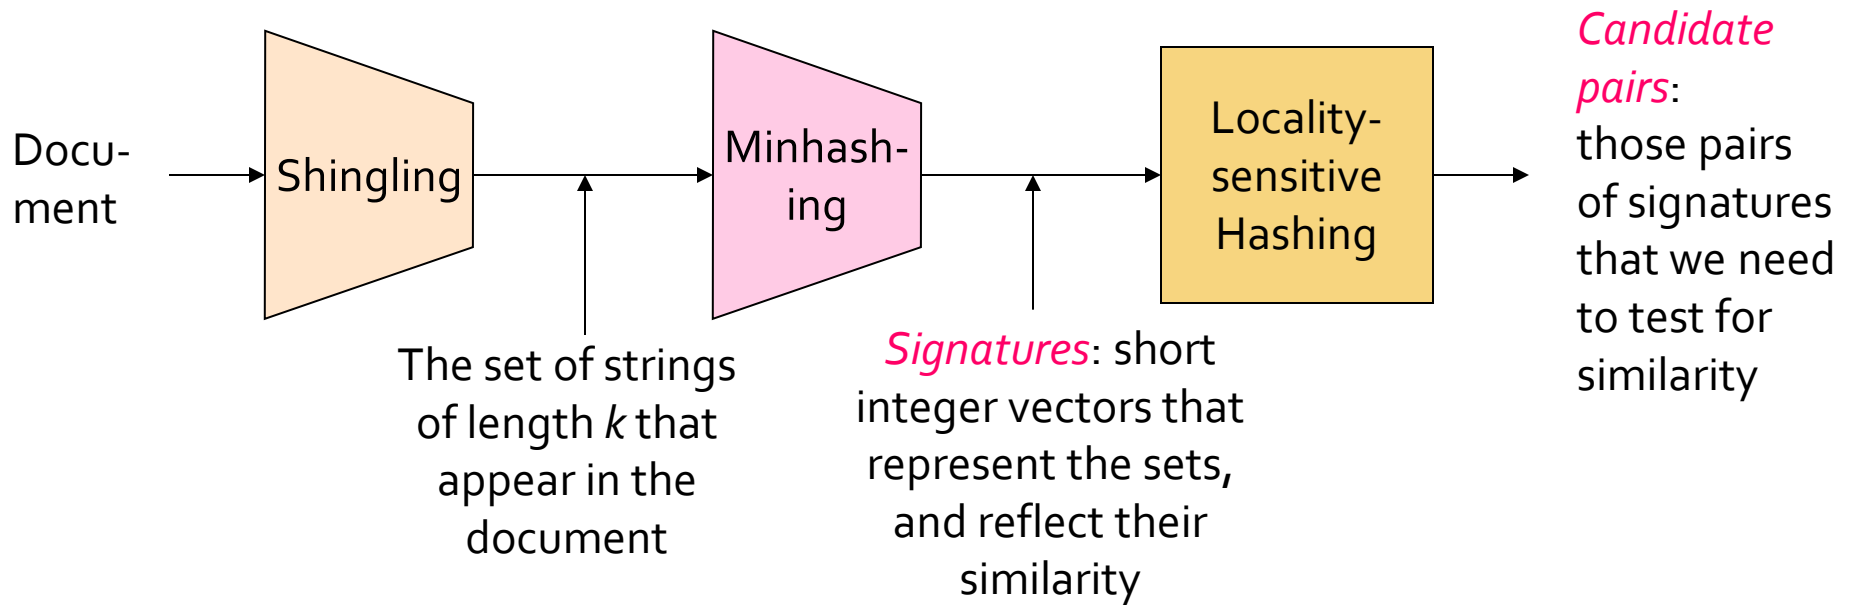
## Data is infinite:

- Mining data streams
- Advertising on the Web

## Applications:

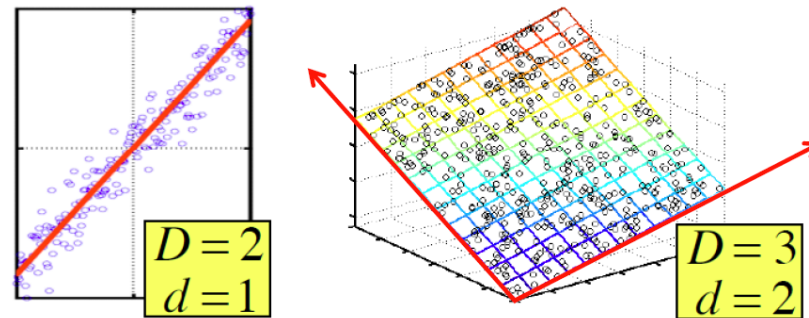
- Association Rules
- Recommender systems

# (1) Finding “similar” sets

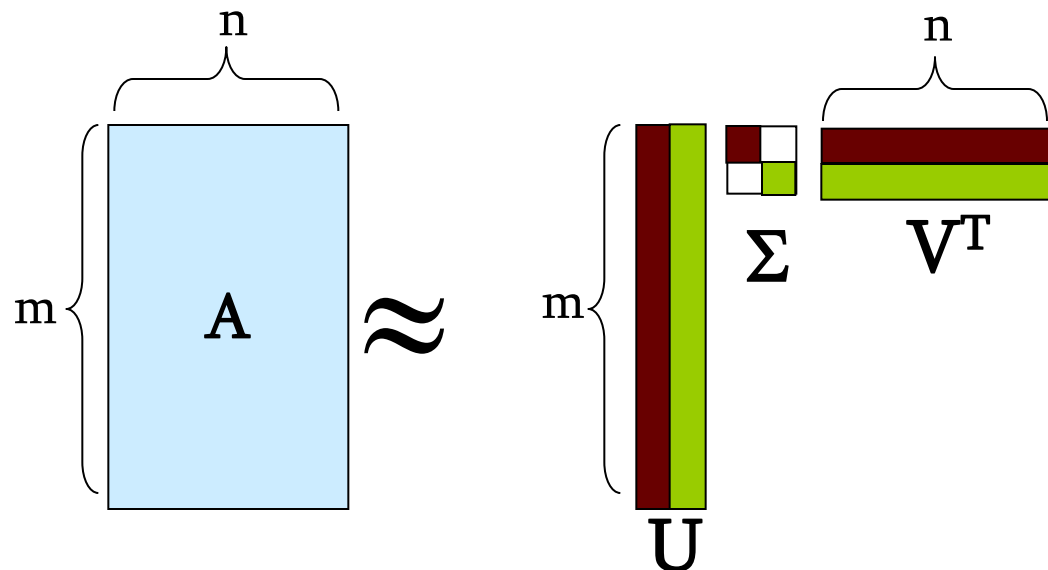


1. **Shingling**: Convert docs to sets
2. **Minhashing**: Convert large sets to short signatures, while preserving similarity
3. **Locality-sensitive hashing**: Focus on pairs of signatures likely to be of similar documents

## (2) Dimensionality Reduction



$$\mathbf{A} \approx \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T = \sum_i \sigma_i \mathbf{u}_i \circ \mathbf{v}_i$$

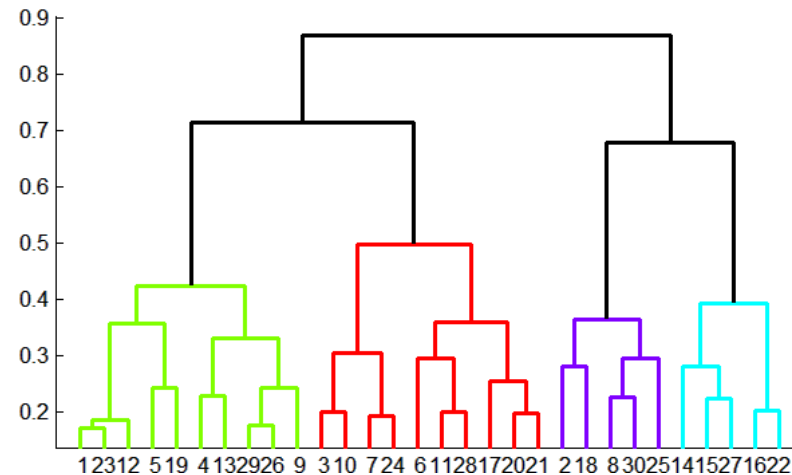


# (3) Clustering

## ■ Hierarchical:

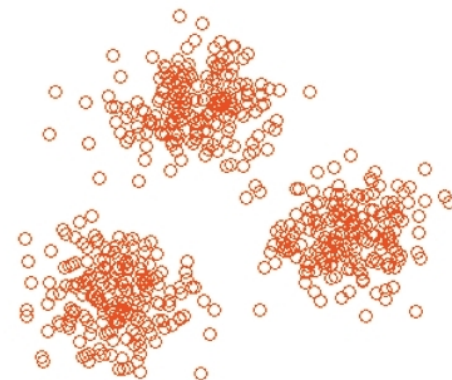
### ■ **Agglomerative** (bottom up):

- Initially, each point is a cluster
- Repeatedly combine the two “nearest” clusters into one
- Represent a cluster by its **centroid** or **clustroid**



## ■ **Point Assignment:**

- Maintain a set of clusters
- Points belong to “nearest” cluster



# High-dim data methods: Comparison

## ■ LSH:

- Find **somewhat** similar pairs of items while avoiding  $O(N^2)$  comparisons

## ■ Clustering:

- Assign points into a **pre-specified** number of **clusters**
  - Each point belongs to a single cluster
  - Summarize the cluster by a centroid

## ■ SVD (dimensionality reduction):

- Want to explore/exploit **correlations** in the data
- Some dimensions may be irrelevant
- Useful for visualization, removing noise from the data, detecting anomalies

# When to use which method?

- **Find all similar pairs of items: LSH**
  - Have to know the threshold ahead of time
  - Allow for some error
- **Identify clusters (structure in data): k-means**
  - $k$  is usually relatively small (10~1000)
  - Useful for identifying ‘types’ or ‘classes’ of datapoints
- **Build low-dimensional representation of data: SVD**
  - More robust (noise-free) similarity computation
  - Data compression (memory saving, speed-up)

# How it all fits together?

## Data is high-dimensional:

- Locality Sensitive Hashing
- Dimensionality reduction
- Clustering

## The data is a graph:

- Link Analysis: PageRank, TrustRank, Hubs & Authorities

## Data is labeled (Machine Learning):

- kNN, Perceptron, SVM, Decision Trees

## Data is infinite:

- Mining data streams
- Advertising on the Web

## Applications:

- Association Rules
- Recommender systems

# Link Analysis: PageRank

- Rank nodes using the network link structure

- PageRank:

- Link voting:

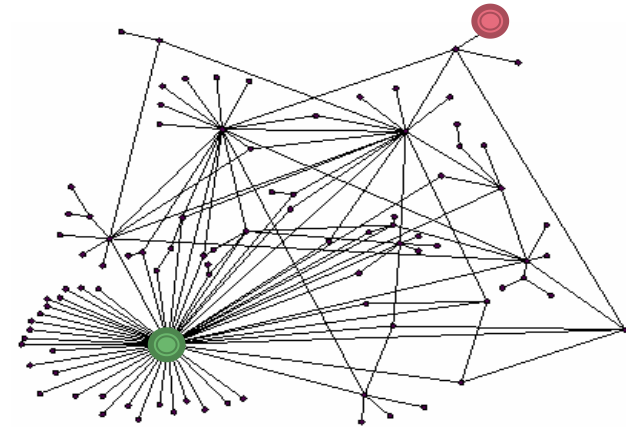
- Page of importance  $x$  has  $n$  out-links, each gets  $x/n$  votes
    - Page  $R$ 's importance is the sum of the votes on its in-links

- Complications: Spider traps, Dead-ends

- Solution: At each step, random surfer has 2 options

- With probability  $\beta$ , follow a link at random
    - With prob.  $1-\beta$ , jump to some page **uniformly** at random

- Power method to compute PageRank





# PPR, SimRank, HITS

## ■ Personalized (topic specific) PageRank

- Random walker teleports to a preselected set of nodes

## ■ Random Walk with Restarts

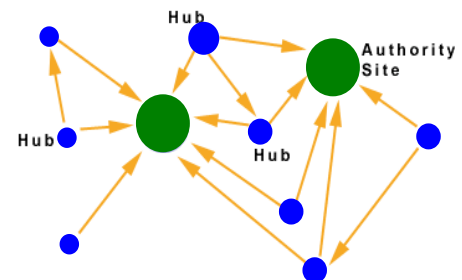
- Random walker always jumps back to the starting node

## ■ SimRank

- **Measure similarity between items**
- **$k$ -partite graph with  $k$  types of nodes**
- Perform a random-walk with restarts from node  **$N$**
- Resulting prob. distrib. is similarity of other nodes to  **$N$**

## ■ Hubs & Authorities

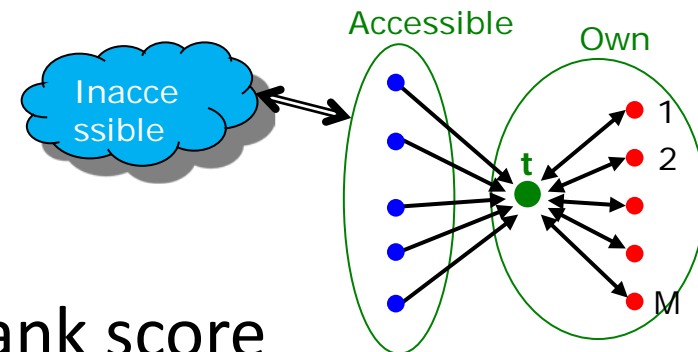
- **Experts vs. Content provides**
- Principle of repeated improvement



# WebSpam and PageRank

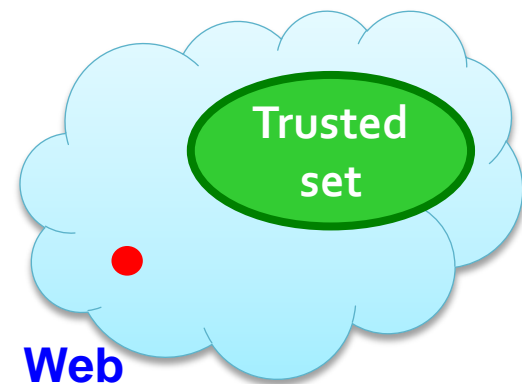
## ■ Web spam farming

- Architecture of a spam farm
- Effect of spam farms on PageRank score



## ■ TrustRank

- Topic specific PageRank with a teleport set of “trusted” pages
- Spam Mass of a page

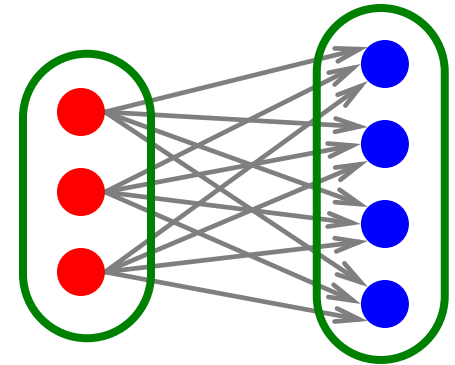


# Analysis of Large Graphs

## ■ Detecting clusters of densely connected nodes

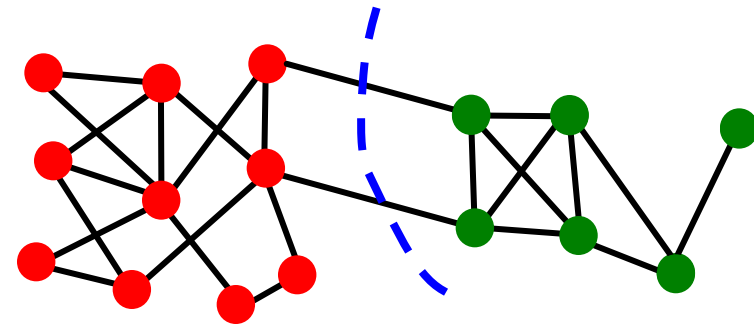
- **Trawling:** Discover complete bipartite subgraphs

- Frequent itemset mining



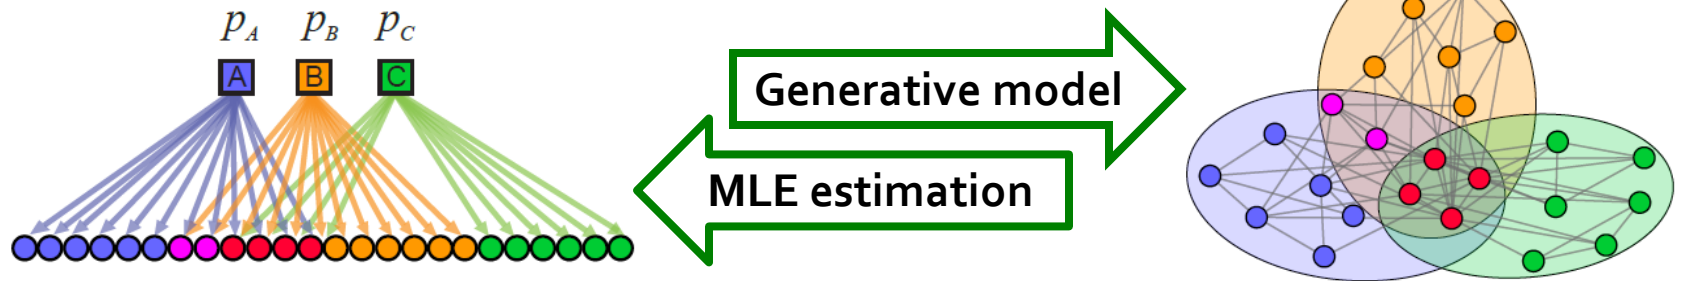
- **Graph partitioning:** “cut” few edges to separate the graph in two pieces

- Conductance
- Computing a sweep
- PageRank-Nibble

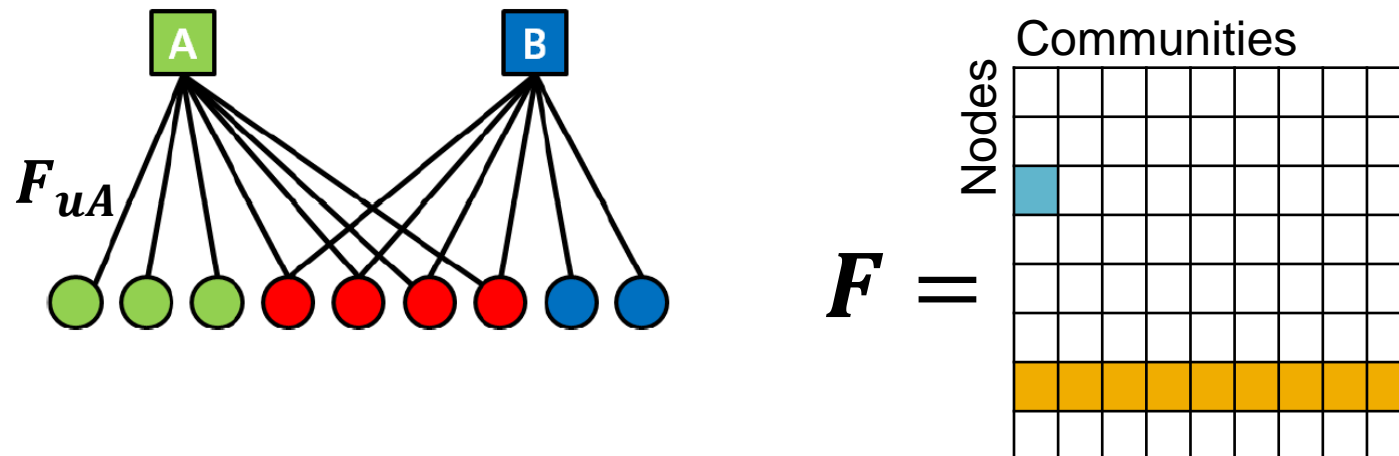


# Analysis of Large Graphs

## ■ AGM (Affiliation Graph Model)



## ■ BigCLAM (CLuster Affiliation Model)



# How it all fits together?

## Data is high-dimensional:

- Locality Sensitive Hashing
- Dimensionality reduction
- Clustering

## The data is a graph:

- Link Analysis: PageRank, TrustRank, Hubs & Authorities

## Data is labeled (Machine Learning):

- kNN, Perceptron, SVM, Decision Trees

## Data is infinite:

- Mining data streams
- Advertising on the Web

## Applications:

- Association Rules
- Recommender systems

# Support Vector Machines

- **Prediction =  $\text{sign}(w \cdot x + b)$**

- Model parameters  $w, b$

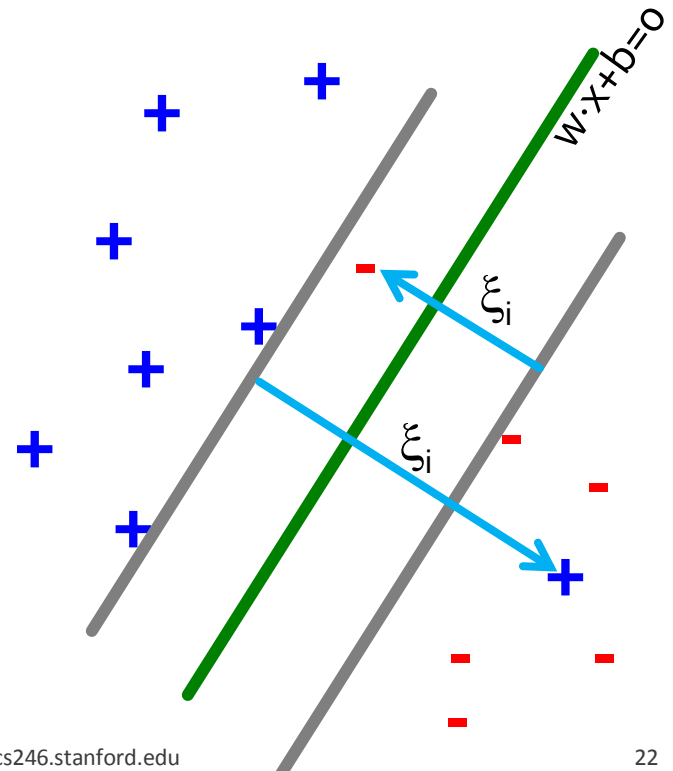
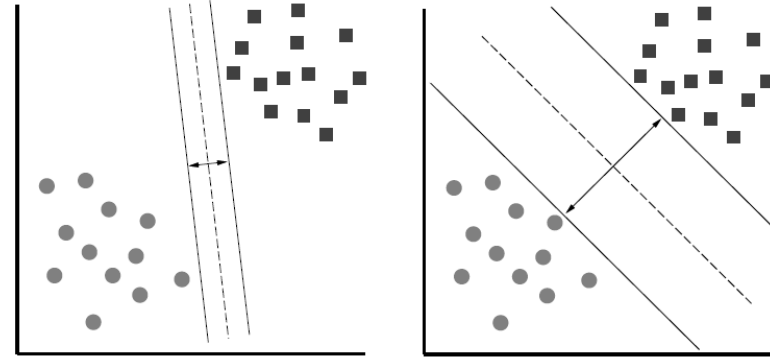
- **Margin:**  $\gamma = \frac{\|w\|}{w \cdot w} = \frac{1}{\|w\|}$

- **SVM optimization problem:**

$$\min_{w, b, \xi_i \geq 0} \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \xi_i$$

$$s.t. \forall i, y_i(w \cdot x_i + b) \geq 1 - \xi_i$$

- Find  $w, b$  using **Stochastic gradient descent**



# Decision Trees: PLANET

## ■ Building decision trees using MapReduce

### ■ How to predict?

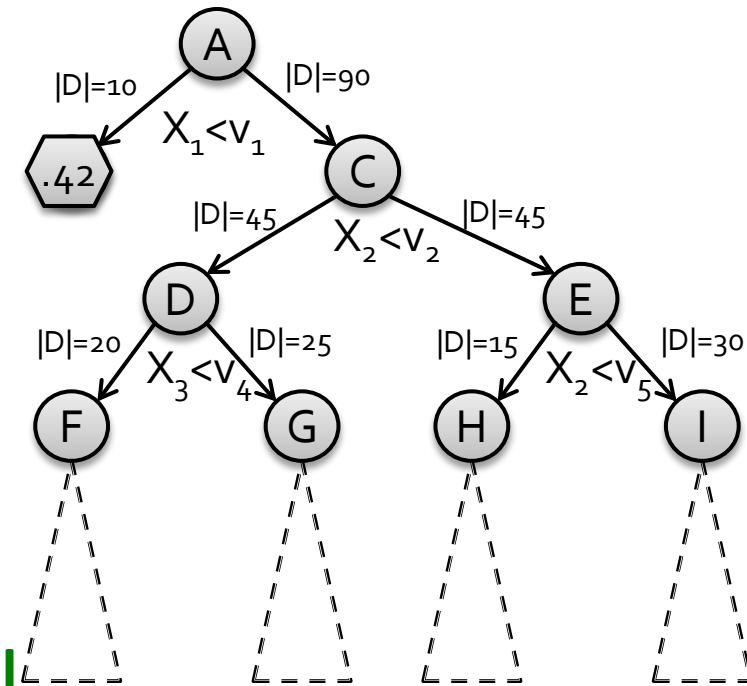
- **Predictor:** avg.  $y_i$  of the examples in the leaf

### ■ When to stop?

- # of examples in the leaf is small

### ■ How to build?

- One MapReduce job per level
  - Need to compute split quality for each attribute and each split value for each current leaf



#### Algorithm 1 FindBestSplit

Require: Node  $n$ , Data  $D \subseteq D^*$

- 1:  $(n \rightarrow \text{split}, D_L, D_R) = \text{FindBestSplit}(D)$
- 2: if StoppingCriteria( $D_L$ ) then
- 3:    $n \rightarrow \text{left\_prediction} = \text{FindPrediction}(D_L)$
- 4: else
- 5:      $\text{FindBestSplit}(n \rightarrow \text{left}, D_L)$
- 6: if StoppingCriteria( $D_R$ ) then
- 7:    $n \rightarrow \text{right\_prediction} = \text{FindPrediction}(D_R)$
- 8: else
- 9:      $\text{FindBestSplit}(n \rightarrow \text{right}, D_R)$

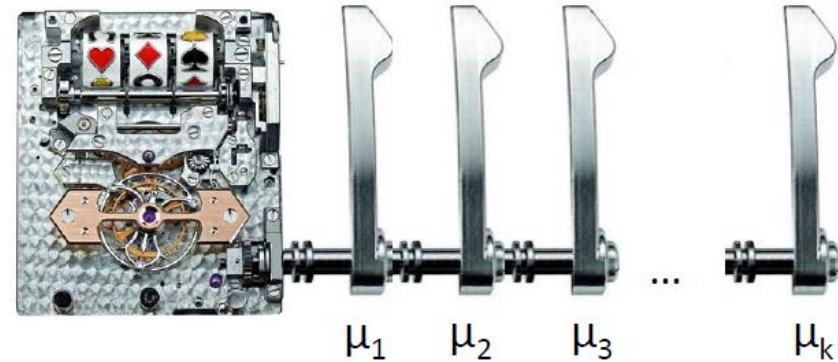
# When to use which method?

- **SVM: Classification**
  - **Millions of numerical features** (e.g., documents)
  - Simple (linear) decision boundary
  - Hard to interpret model
- **k-NN: Classification or regression**
  - (Many) numerical features
  - Many parameters to play with – distance metric,  $k$ , weighting, ... **there is no simple way to set them!**
- **Decision Trees: Classification or Regression**
  - Relatively few features (handles categorical features)
  - Complicated decision boundary: **Overfitting!**
  - **Easy to explain/interpret the classification**
  - **Bagged Decision Trees** – very, very hard to beat!



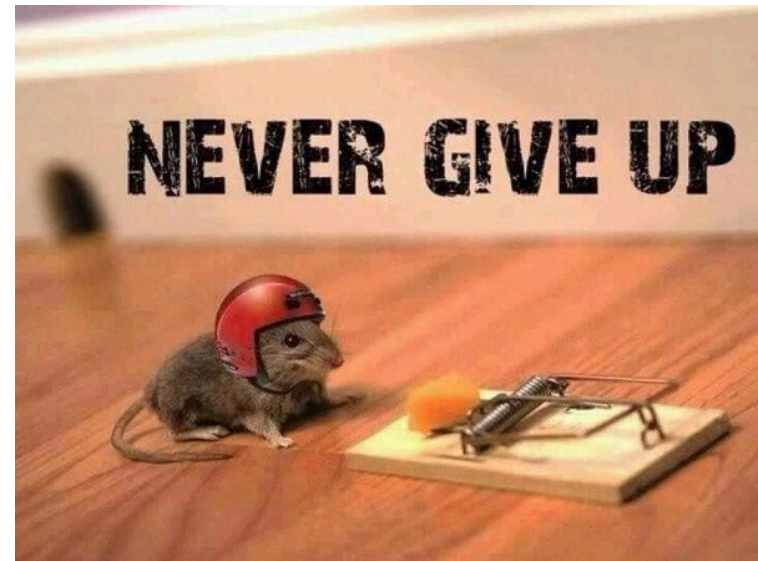
# Learning Through Experimentation

- **Learning through experimentation**
  - Exploration-Exploitation tradeoff
  - Regret
- **Multiarmed Bandits**
  - Epsilon-Greedy
  - UCB1 algorithm
- **Submodular function optimization**
  - Coverage
  - Greedy and Lazy-Greedy algorithms
  - Multiplicative Weights algorithm



# What if “ML alg. doesn’t work”?

- Compare error on the train/test set
- Plot error vs. (regularization) parameter
- Compare performance to a simple baseline
- Build synthetic datasets for which you know your method should work
- **Think about:**
  - the prediction problem
  - Error metrics
  - Model assumptions
  - Properties of the data



# If the ML algorithm doesn't work

- **Get more training data**
  - Sometimes more data doesn't help but often it does
- **Try a smaller set a features**
  - Carefully select small subset
  - You can do this by hand, or use SVD
- **Try getting additional features**
  - **LOOK** at the data
  - Can be very time consuming
- **Adding polynomial features**
  - Include  $\mathbf{x}$  and  $\mathbf{x}^2$  as features
- **Building your own, new, better features**
  - Based on your knowledge of the problem
- **Try decreasing or increasing C**
  - Change how important the regularization term is

# How it all fits together?

## Data is high-dimensional:

- Locality Sensitive Hashing
- Dimensionality reduction
- Clustering

## The data is a graph:

- Link Analysis: PageRank, TrustRank, Hubs & Authorities

## Data is labeled (Machine Learning):

- kNN, Perceptron, SVM, Decision Trees

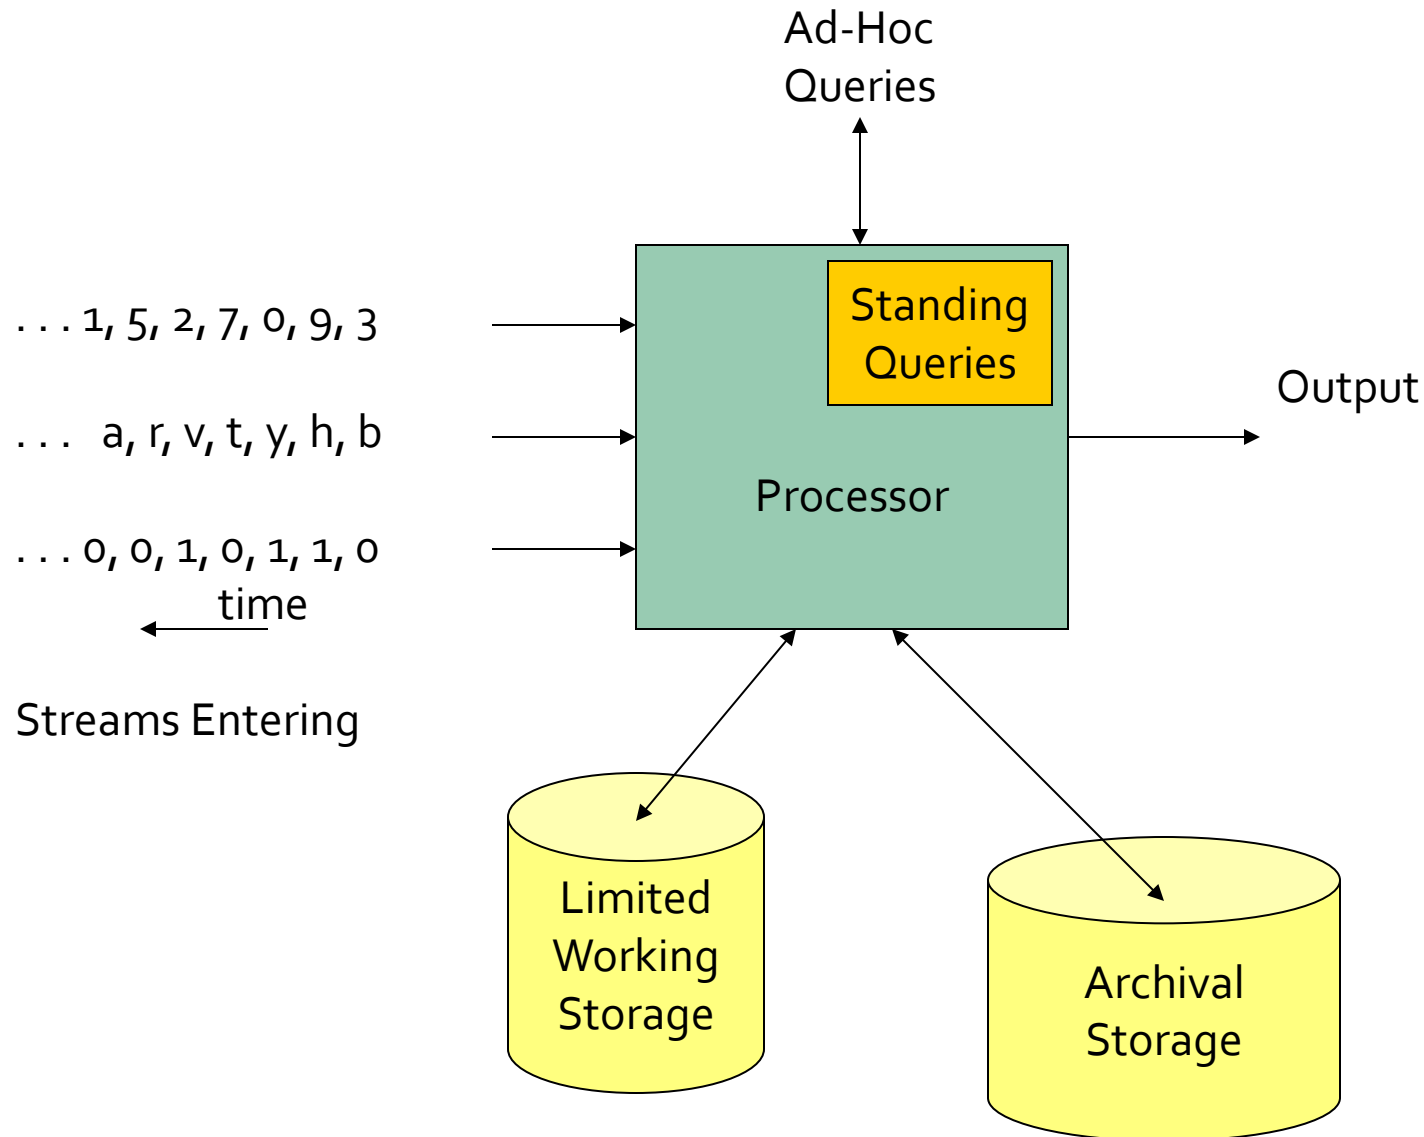
## Data is infinite:

- Mining data streams
- Advertising on the Web

## Applications:

- Association Rules
- Recommender systems

# Mining Data Streams



# Problems on data streams

- **Sampling data from a stream:**

- Each element is included with prob.  $k/N$

- **Queries over sliding windows:**

How many **1**s are in last  $k$  bits?

1001010110001011010101010101011010101010101110101010111010101011101010100010110010

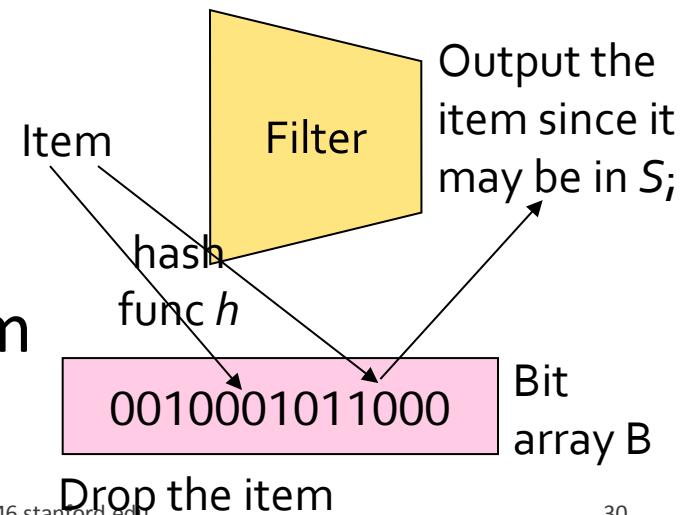
- **Filtering a stream: Bloom filters**

- Filter elements with property  $x$

- **Counting distinct elements:**

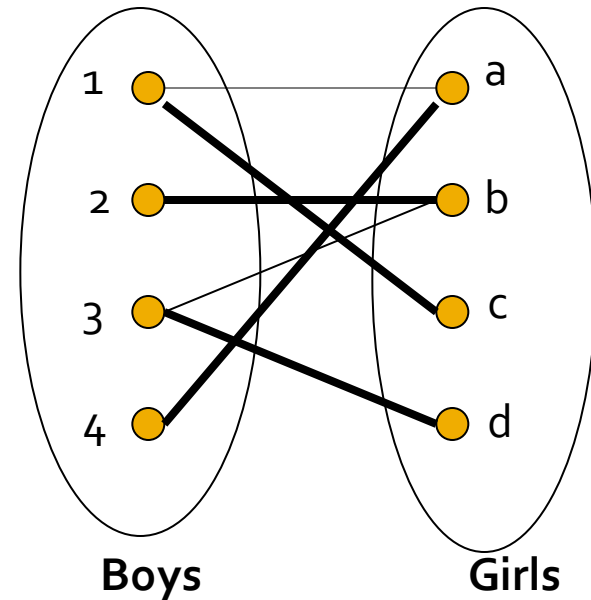
- Number of distinct elements in the last  $k$  elements of the stream

- **Estimating moments**



# Online algorithms & Advertising

- You get to see one input piece at a time, and need to make irrevocable decisions
- **Competitive ratio** =  $\min_{\text{all inputs}} (|M_{\text{my\_alg}}| / |M_{\text{opt}}|)$
- **Adwords problem:**
  - Query arrives to a search engine
  - Several advertisers bid on the query
  - Pick a subset of advertisers whose ads are shown
- **Greedy online matching: competitive ratio  $\geq 1/2$**



# How it all fits together?

## Data is high-dimensional:

- Locality Sensitive Hashing
- Dimensionality reduction
- Clustering

## The data is a graph:

- Link Analysis: PageRank, TrustRank, Hubs & Authorities

## Data is labeled (Machine Learning):

- kNN, Perceptron, SVM, Decision Trees

## Data is infinite:

- Mining data streams
- Advertising on the Web

## Applications:

- Association Rules
- Recommender systems



# Association Rule Discovery

## Market-basket model:

- **Goal:** To identify items that are bought together by sufficiently many customers
- **Approach:** Process the sales data collected with barcode scanners to find dependencies among items

<i>TID</i>	<i>Items</i>
1	Bread, Coke, Milk
2	Beer, Bread
3	Beer, Coke, Diaper, Milk
4	Beer, Bread, Diaper, Milk
5	Coke, Diaper, Milk

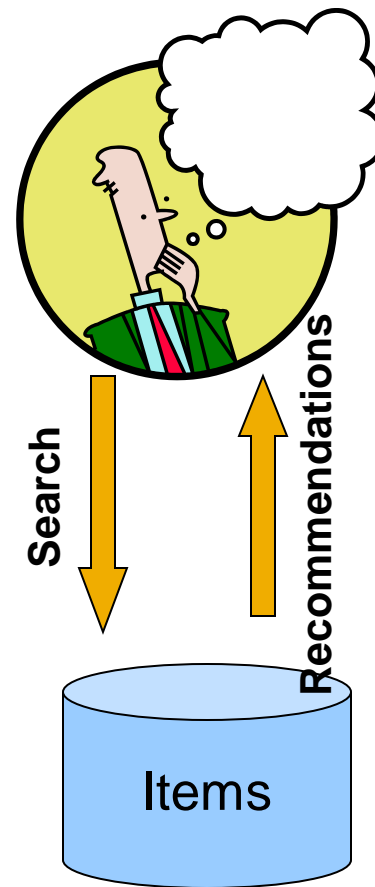
Rules Discovered:

$\{\text{Milk}\} \rightarrow \{\text{Coke}\}$

$\{\text{Diaper, Milk}\} \rightarrow \{\text{Beer}\}$

# Recommender Systems

- **User-user collaborative filtering**
  - Consider user  $c$
  - Find set  $D$  of other users whose ratings are “similar” to  $c$ ’s ratings
  - Estimate user’s ratings based on the ratings of users in  $D$
- **Item-item collaborative filtering**
  - Estimate rating for item based on ratings for similar items
- **Profile based**



# Latent Factor Models: Netflix

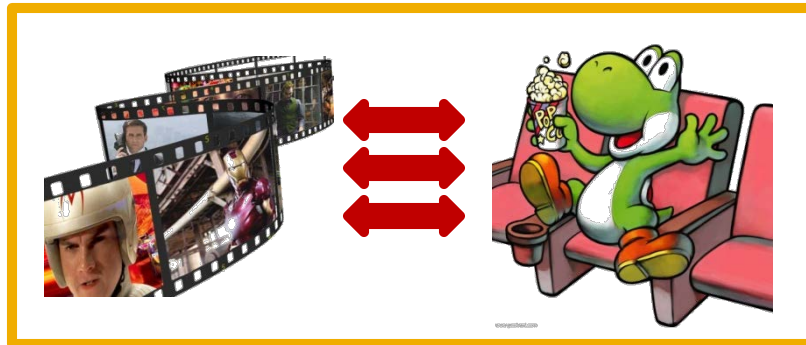
user bias



movie bias



user-movie interaction



## Baseline predictor

- Separates users and movies
- Benefits from insights into user's behavior

## User-Movie interaction

- Characterizes the matching between users and movies
- Attracts most research in the field

$$\min_{Q, P} \sum_{(u, i) \in R} \left( r_{ui} - (\mu + b_u + b_i + q_i^T p_u) \right)^2$$
$$+ \lambda \left( \sum_i \|q_i\|^2 + \sum_u \|p_u\|^2 + \sum_u \|b_u\|^2 + \sum_i \|b_i\|^2 \right)$$

# When to use which method?

- **Lots of rating data: CF**
  - Easy to tweak, easy to add lots of features/signals
  - Use optimization to learn weights on how to combine features
- **Lots<sup>2</sup> of rating data: CF + Latent factors**
  - Many ratings per user, many ratings per item
  - Depending on the amount of data make the model more/less complex (more/less parameters)
- **Cold start, little data: Profile based**
  - Need to have good user/item features and similarity metric

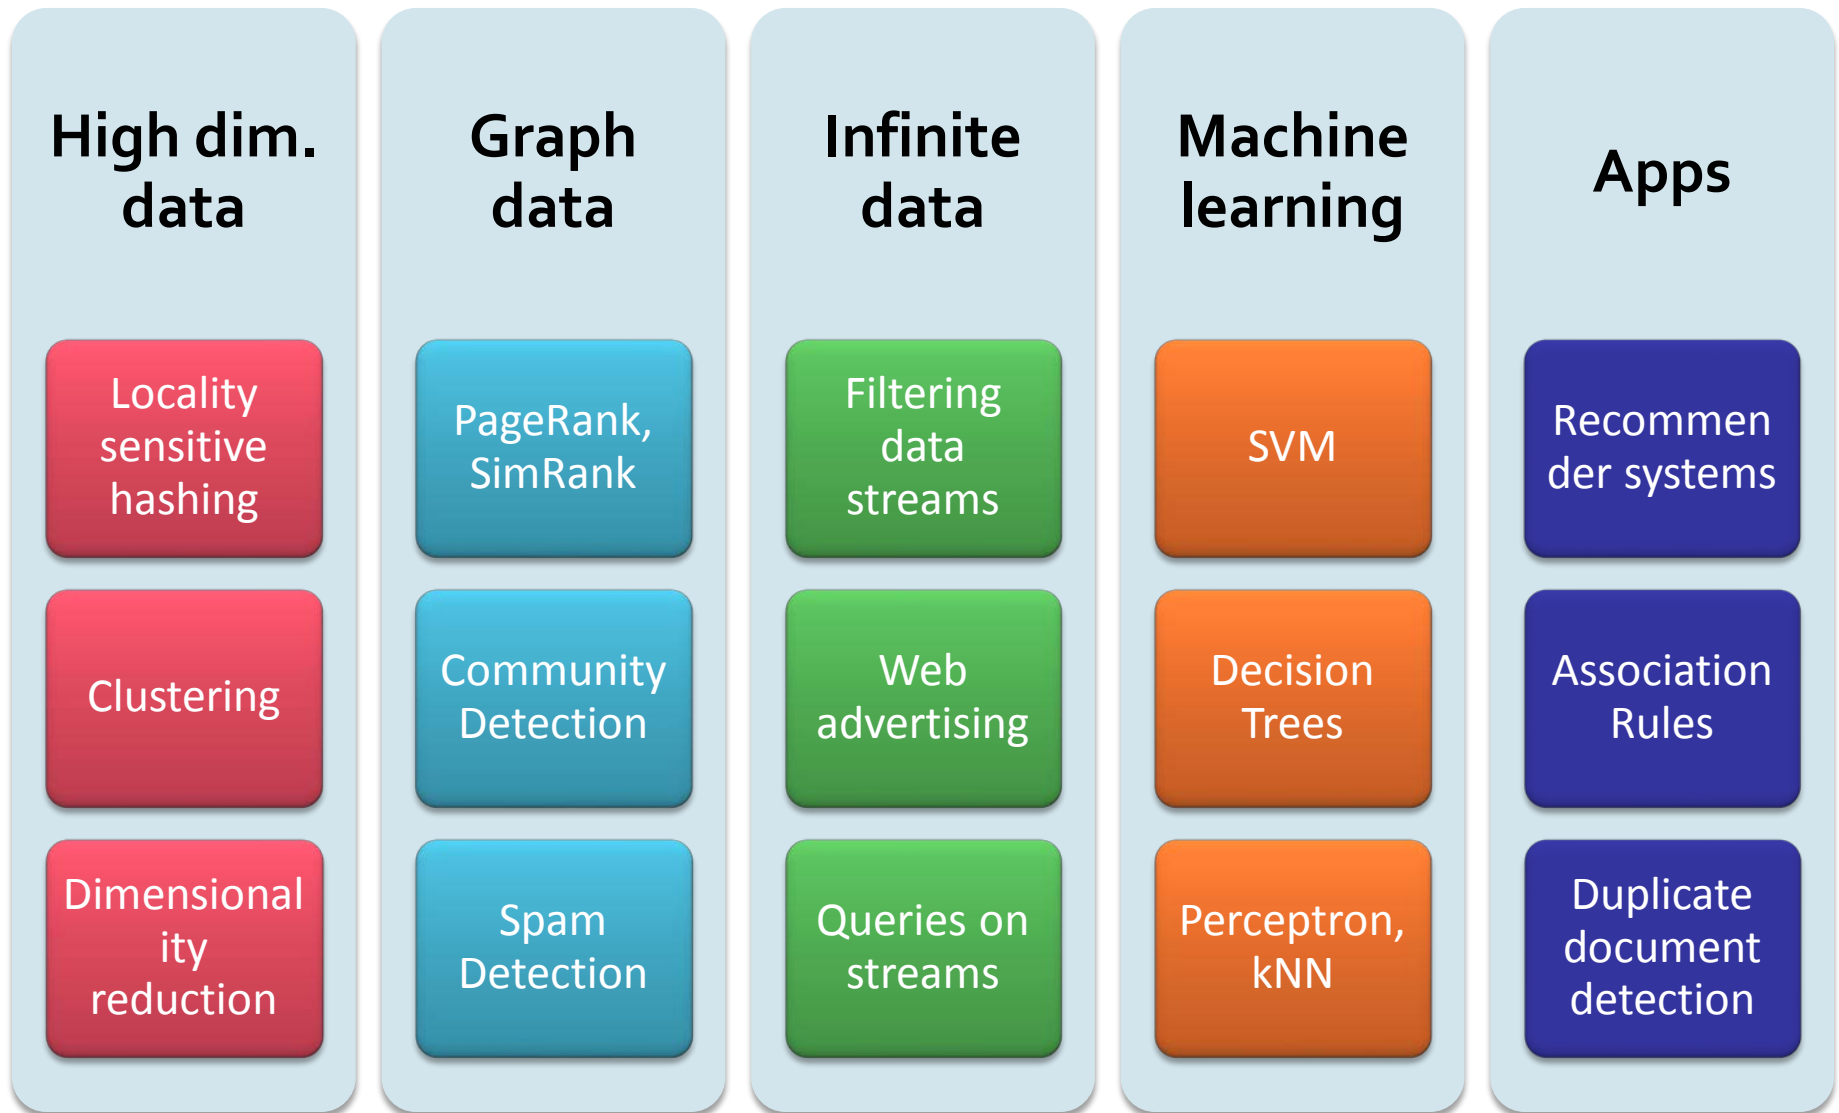
**In closing...**

# What we've learned this quarter

- MapReduce
- Association Rules
- Apriori algorithm
- Finding Similar Items
- Locality Sensitive Hashing
- Random Hyperplanes
- Dimensionality Reduction
- Singular Value Decomposition
- CUR method
- Clustering
- Recommender systems
- Collaborative filtering
- PageRank and TrustRank
- Hubs & Authorities
- k-Nearest Neighbors
- Perceptron
- Support Vector Machines
- Stochastic Gradient Descent
- Decision Trees
- Mining data streams
- Bloom Filters
- Flajolet-Martin
- Advertising on the Web



# Map of Superpowers





# Applying Your Superpowers





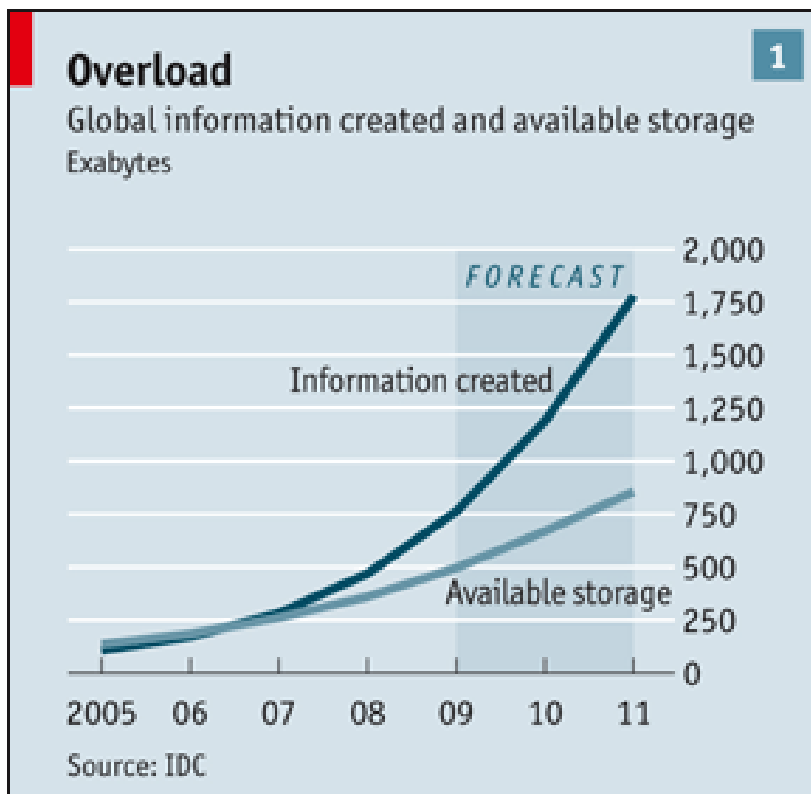
# Applying Your Superpowers



How do you want that data?

# Applying Your Superpowers

- We are producing more data than we are able to store!



[The economist, 2010]



# THE BIG PICTURE

- How to analyze large datasets to discover **models** and **patterns** that are:
  - **Valid:** Hold on new data with some certainty
  - **Novel:** Non-obvious to the system
  - **Useful:** Should be possible to act on the item
  - **Understandable:** Humans should be able to interpret the pattern
- **How to do this using massive data** (that does not fit into main memory)

# What next? Seminars

## ■ Seminars:

- InfoSeminar: <http://i.stanford.edu/infoseminar>
- RAIN Seminar: <http://rain.stanford.edu>

## ■ Conferences:

- **KDD**: ACM Conf. on Knowledge Discovery & Data Mining
- **WSDM**: ACM Conf. on Web Search and Data Mining
- **ICDM**: IEEE International Conf. on Data Mining
- **WWW**: World Wide Web Conference
- **ICML**: International Conf. on Machine Learning
- **VLDB**: Very Large Data Bases

# CS341: Project in Data Mining

- **Data mining research project on real data**
  - Groups of 3 students
  - **We provide interesting data, computing resources (Amazon EC2) and mentoring**
  - **You provide project ideas**
  - There are (practically) no lectures, only individual group mentoring

**Information session:**  
**Tuesday 3/18 7pm in Gates 104**  
(there will be pizza)

# What Next? Courses

- **Other relevant courses**
  - **CS224W**: Social and Information Network Analysis
  - **CS276**: Information Retrieval and Web Search
  - **CS229**: Machine Learning
  - **CS245**: Database System Principles
  - **CS347**: Distributed Databases
  - **CS448g**: Interactive Data Analysis

# What Next? Final Exam

